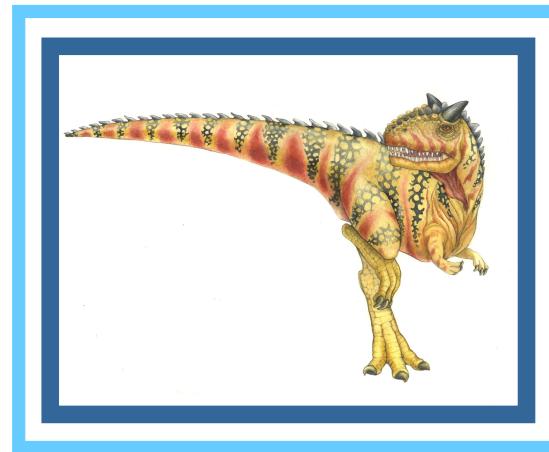


Chapter 9: Main Memory





Objectives

- To provide a detailed description of various ways of organizing memory hardware
- To discuss various memory-management techniques, including paging and segmentation



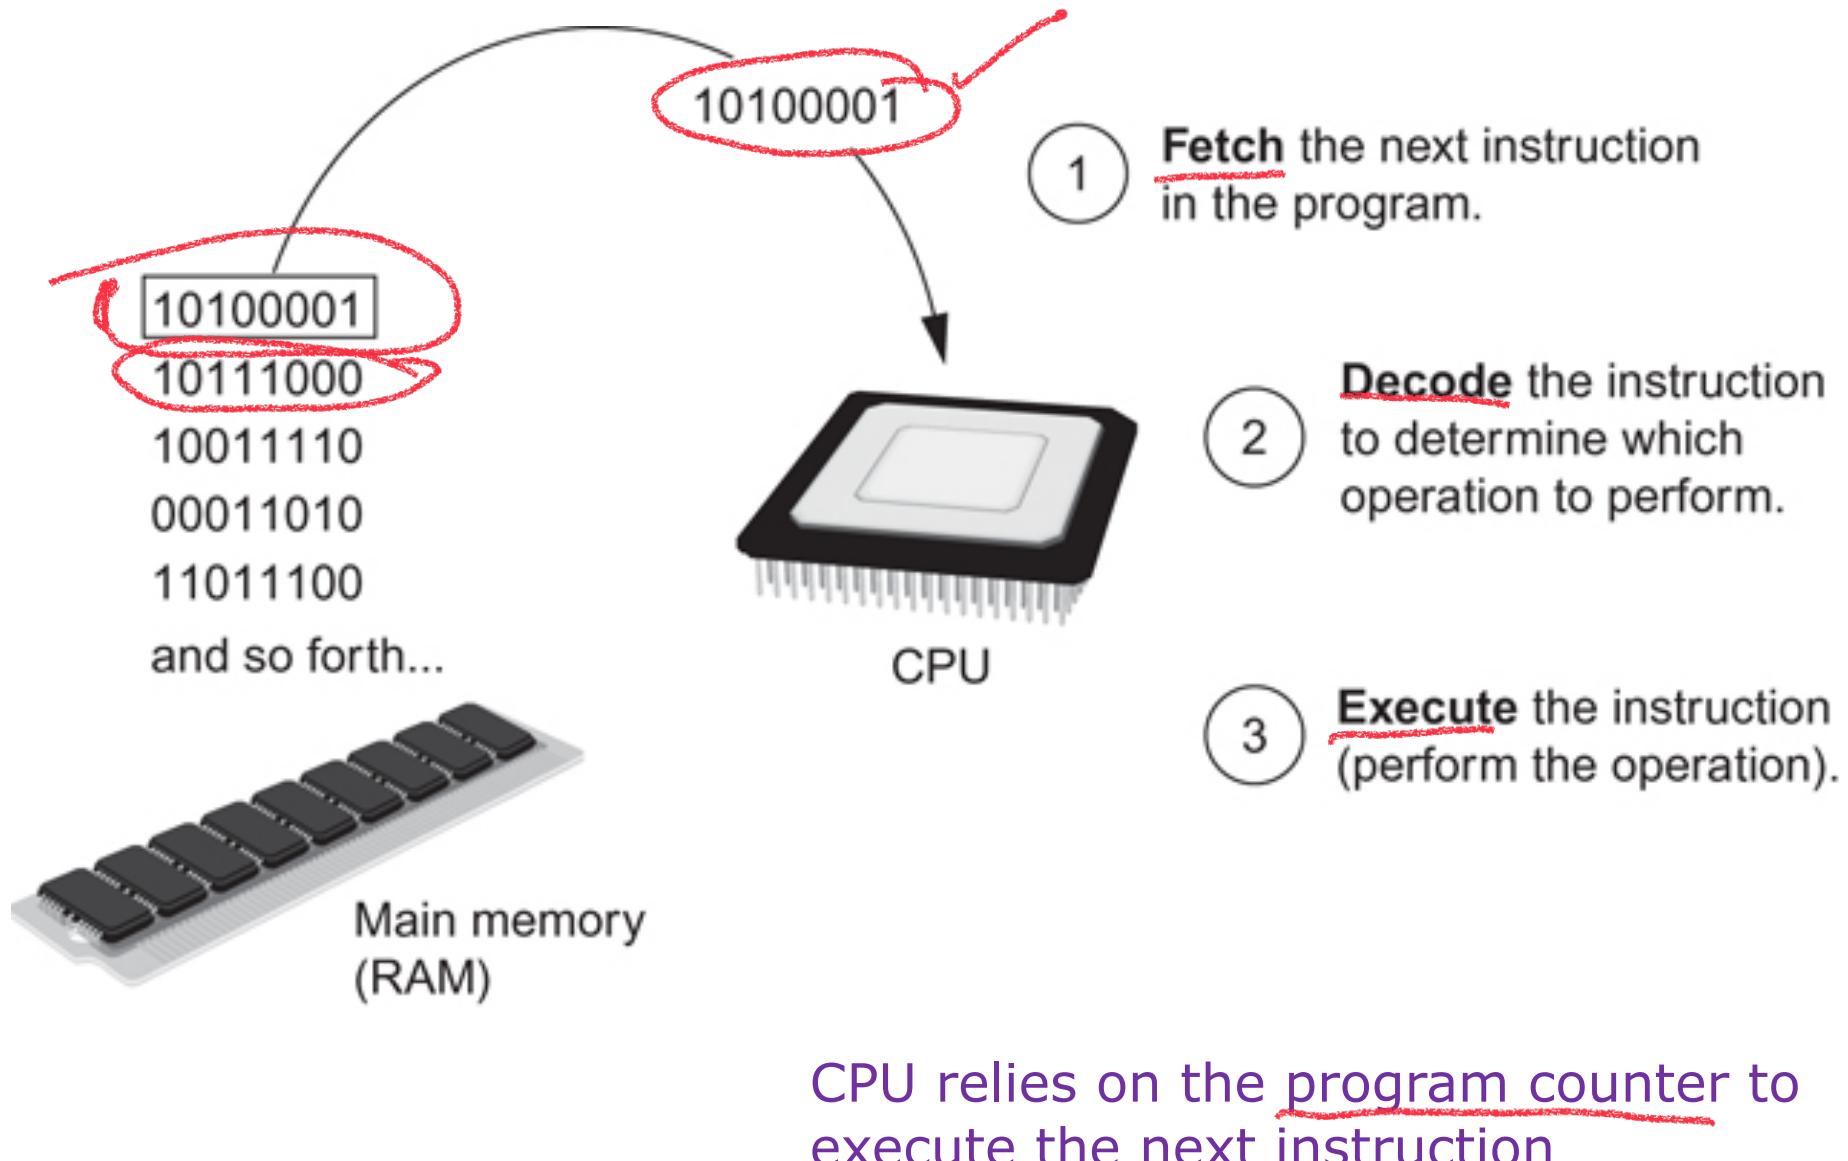


Background

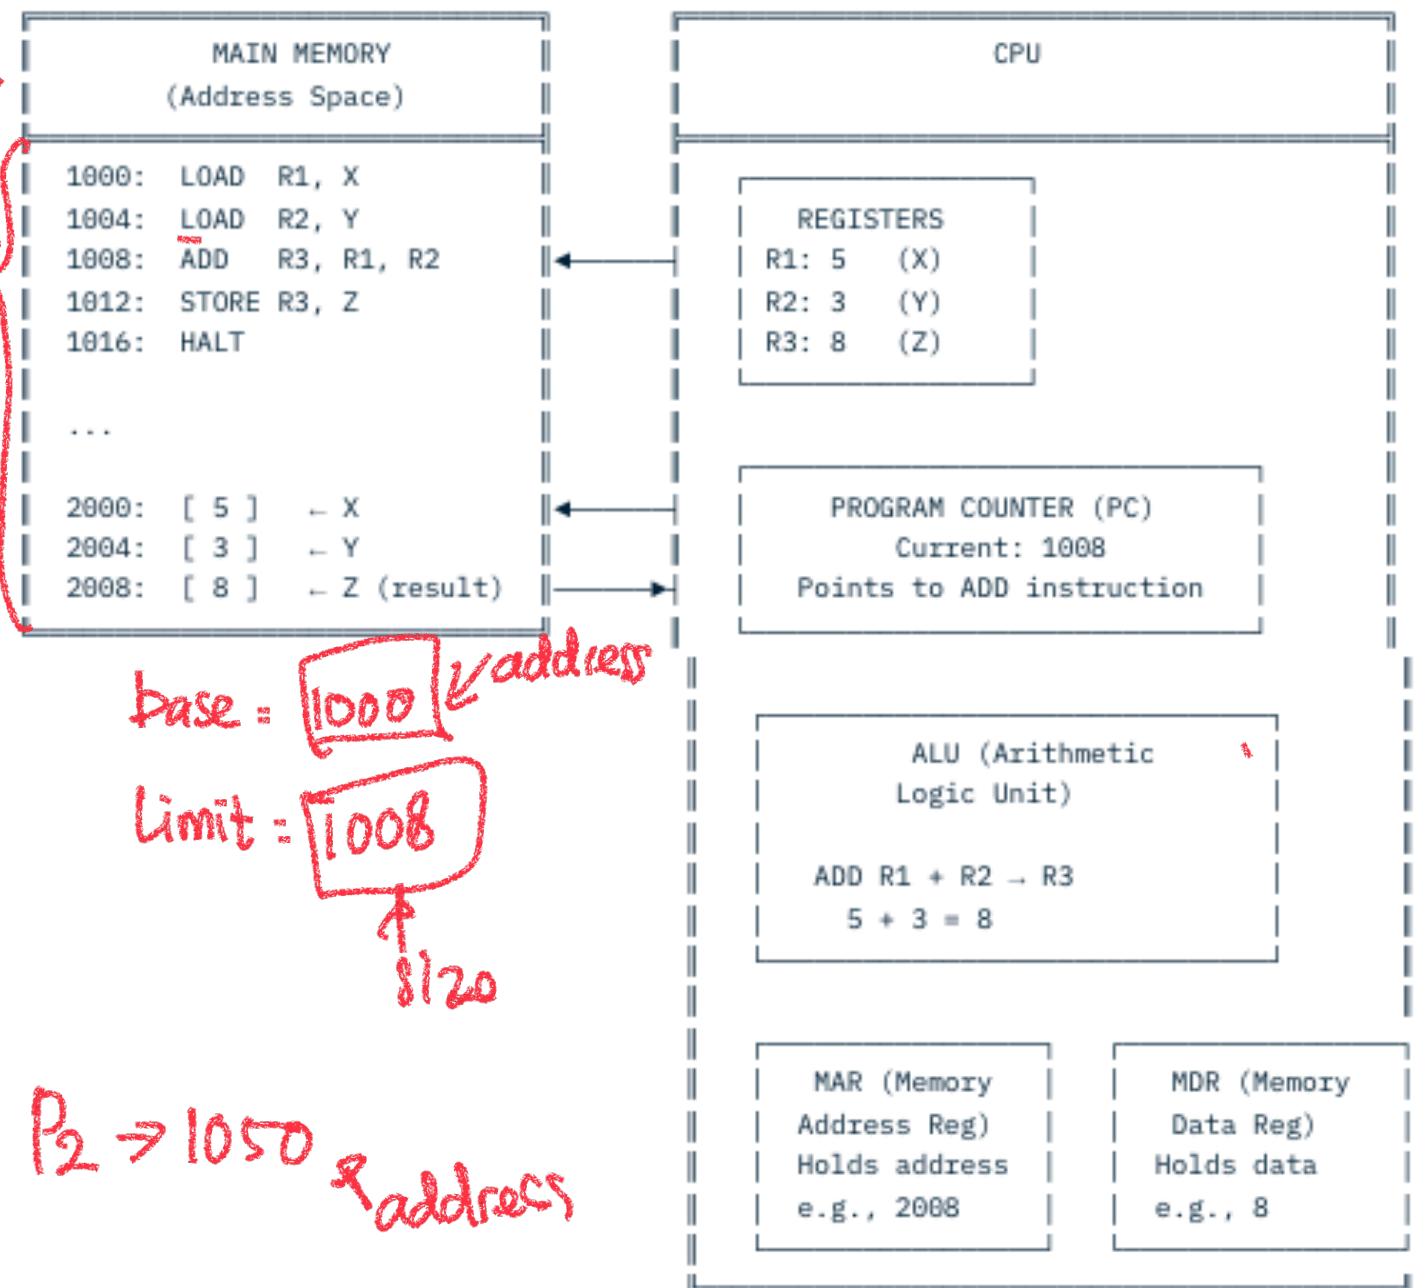
- Program must be brought (from disk) into memory and placed within a process for it to be run
- Main memory and registers are only storage CPU can access directly
- Memory unit only sees a stream of addresses + read requests, or address + data and write requests
- Register access in one CPU clock (or less)
 - Main memory can take many cycles, causing a **stall**
 - **Cache** sits between main memory and CPU registers



Recall : How a Program Works



Simple Program: $Z = X + Y$





Background

- Protection of memory required to ensure correct operation
 - Make sure that each process has a separate memory space
 - fundamental to having multiple processes loaded in memory for concurrent execution

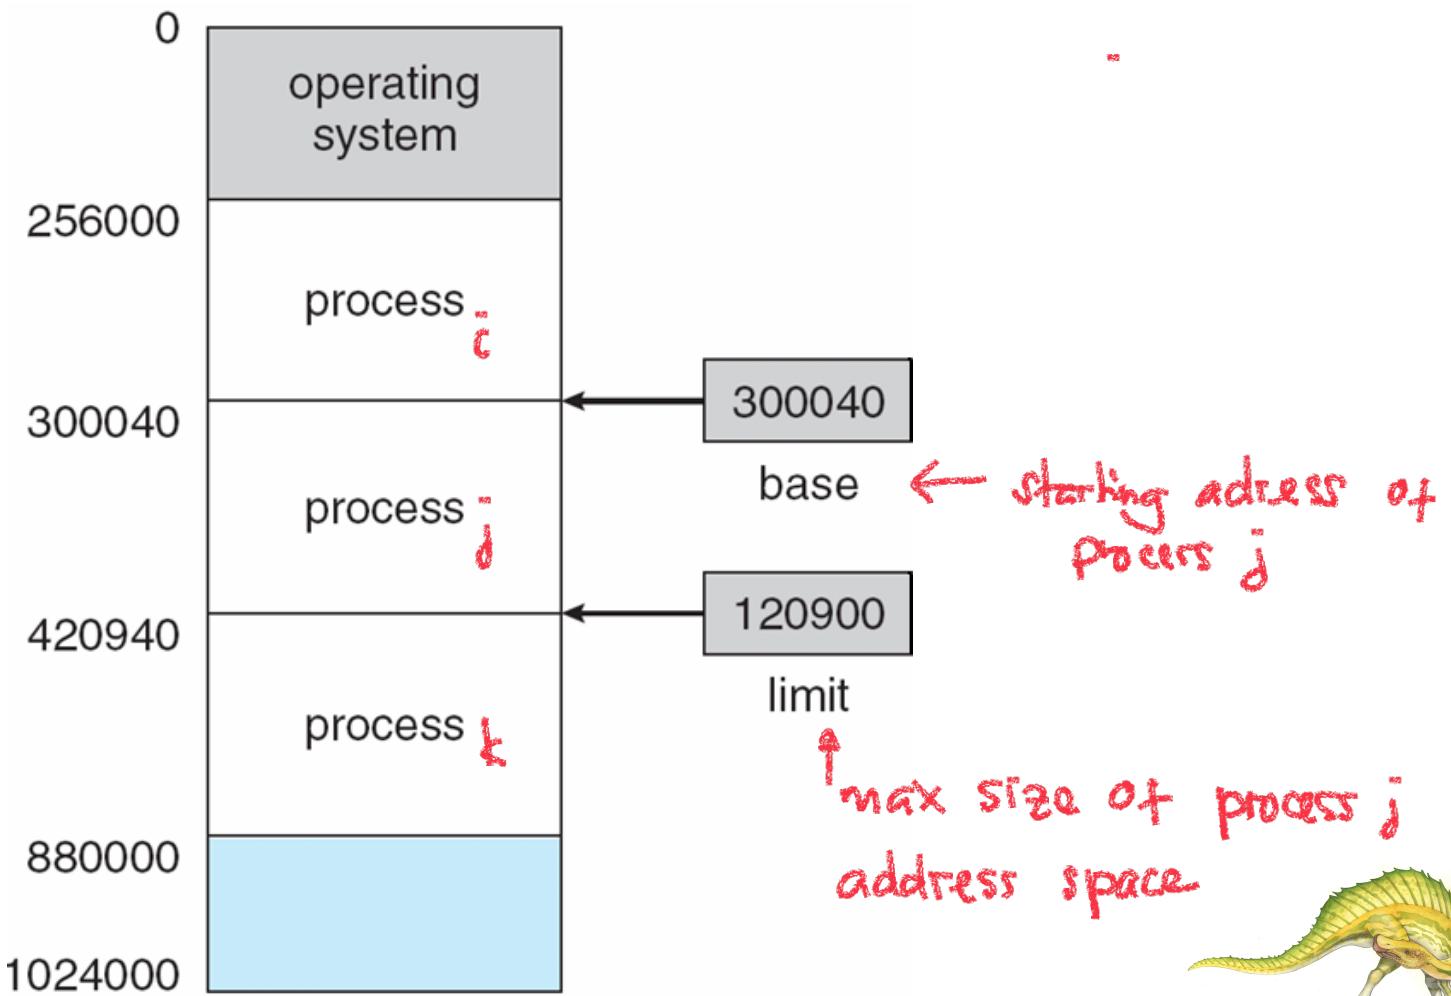
→ main objective of OS¹ memory management





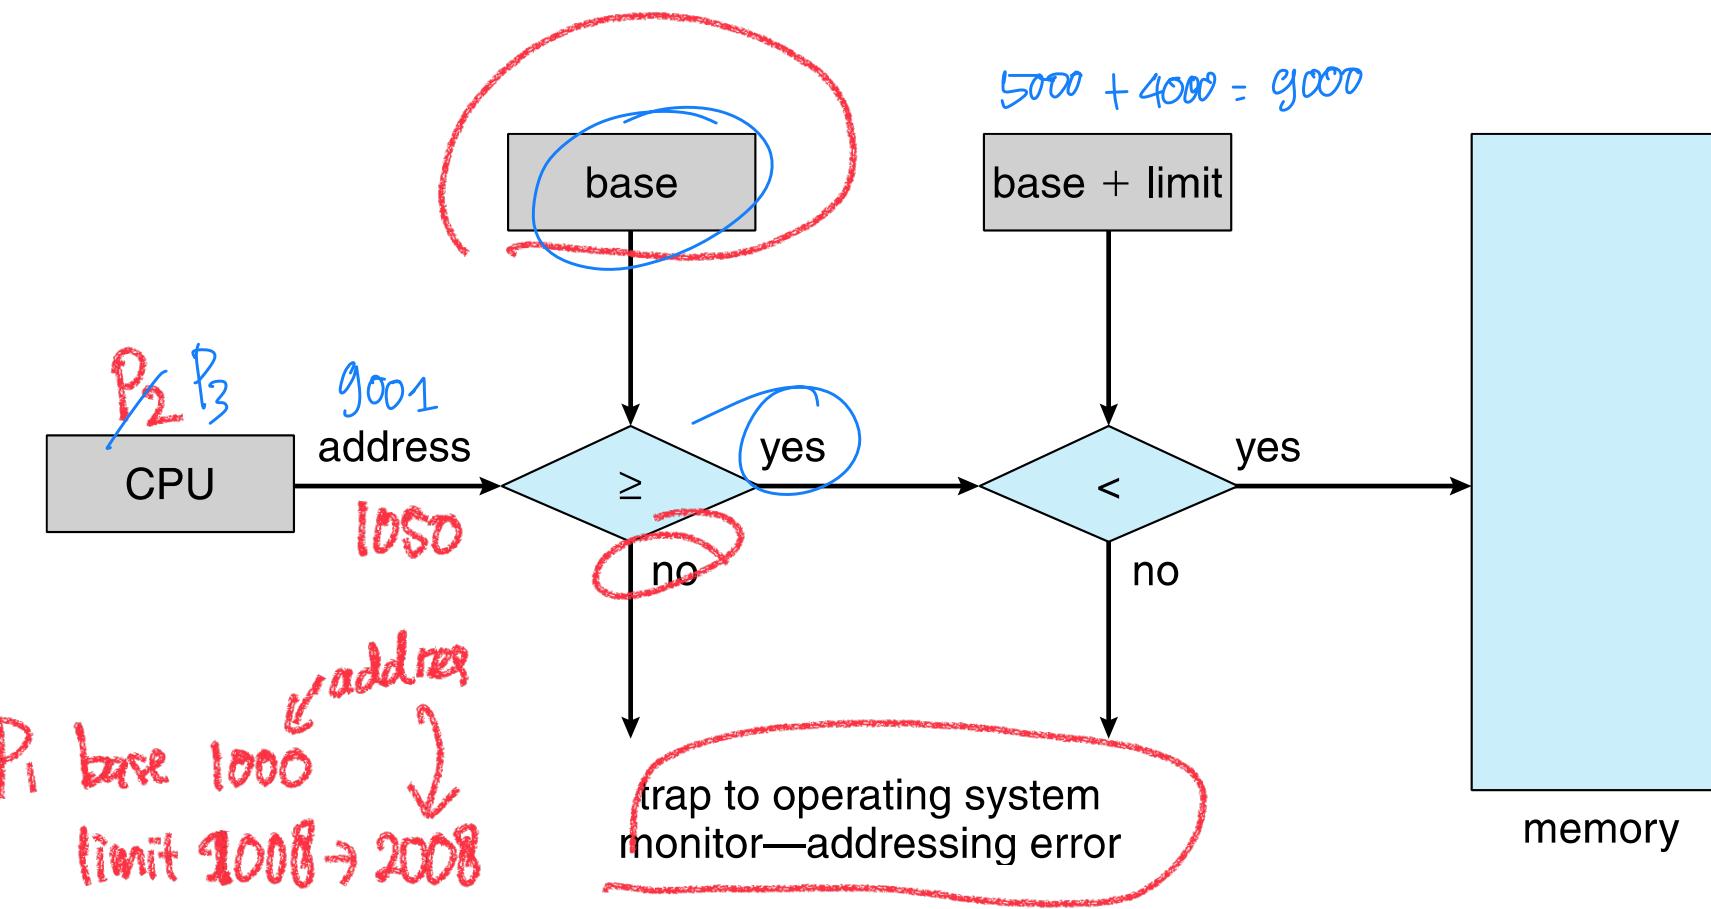
Base and Limit Registers

- A pair of **base** and **limit registers** define the logical address space
- CPU must check every memory access generated in user mode to be sure it is between base and limit for that user





Hardware Address Protection



P_2 base 2009
limit 1008

P_2 want 1050 address

Simple Scheme

P_3 : base 5000
limit 4000
size 9001





Address Binding (1)

- Programs on disk, ready to be brought into memory to execute from an **input queue**
 - Without support, must be loaded into address 0000
 - Inconvenient to have first user process physical address always at 0000
 - How can it not be?

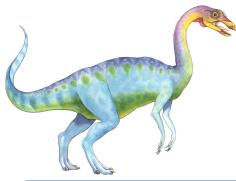




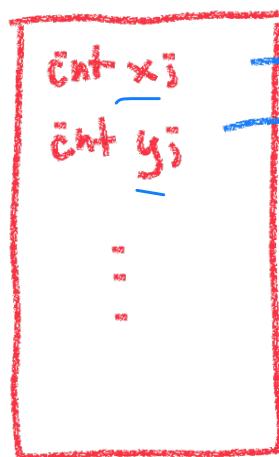
Address Binding (2)

- Further, addresses represented in different ways at different stages of a program's life
 - Source code addresses usually **symbolic** (e.g, a variable might be referred to by its name (`x` or `total`) rather than an exact memory address)
 - Compiled code addresses **bind** to relocatable addresses
 - the compiler translates symbolic addresses into *relocatable addresses*.
 - i.e. “14 bytes from beginning of this module”
 - Linker or loader will bind **relocatable addresses** to **absolute addresses**
 - i.e. 74014
 - Each binding maps one address space to another



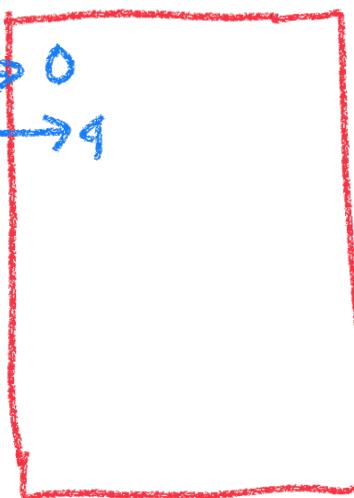


Source File



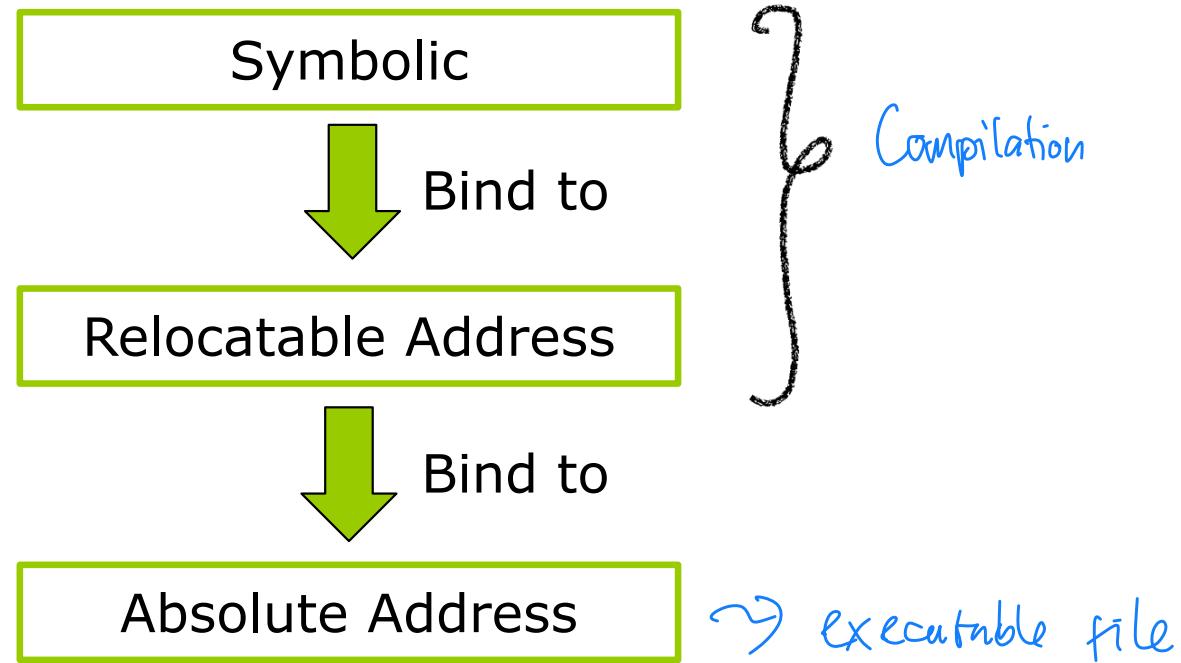
Compiler

Machine Code \rightarrow No variables
yet address





Address Binding (3)





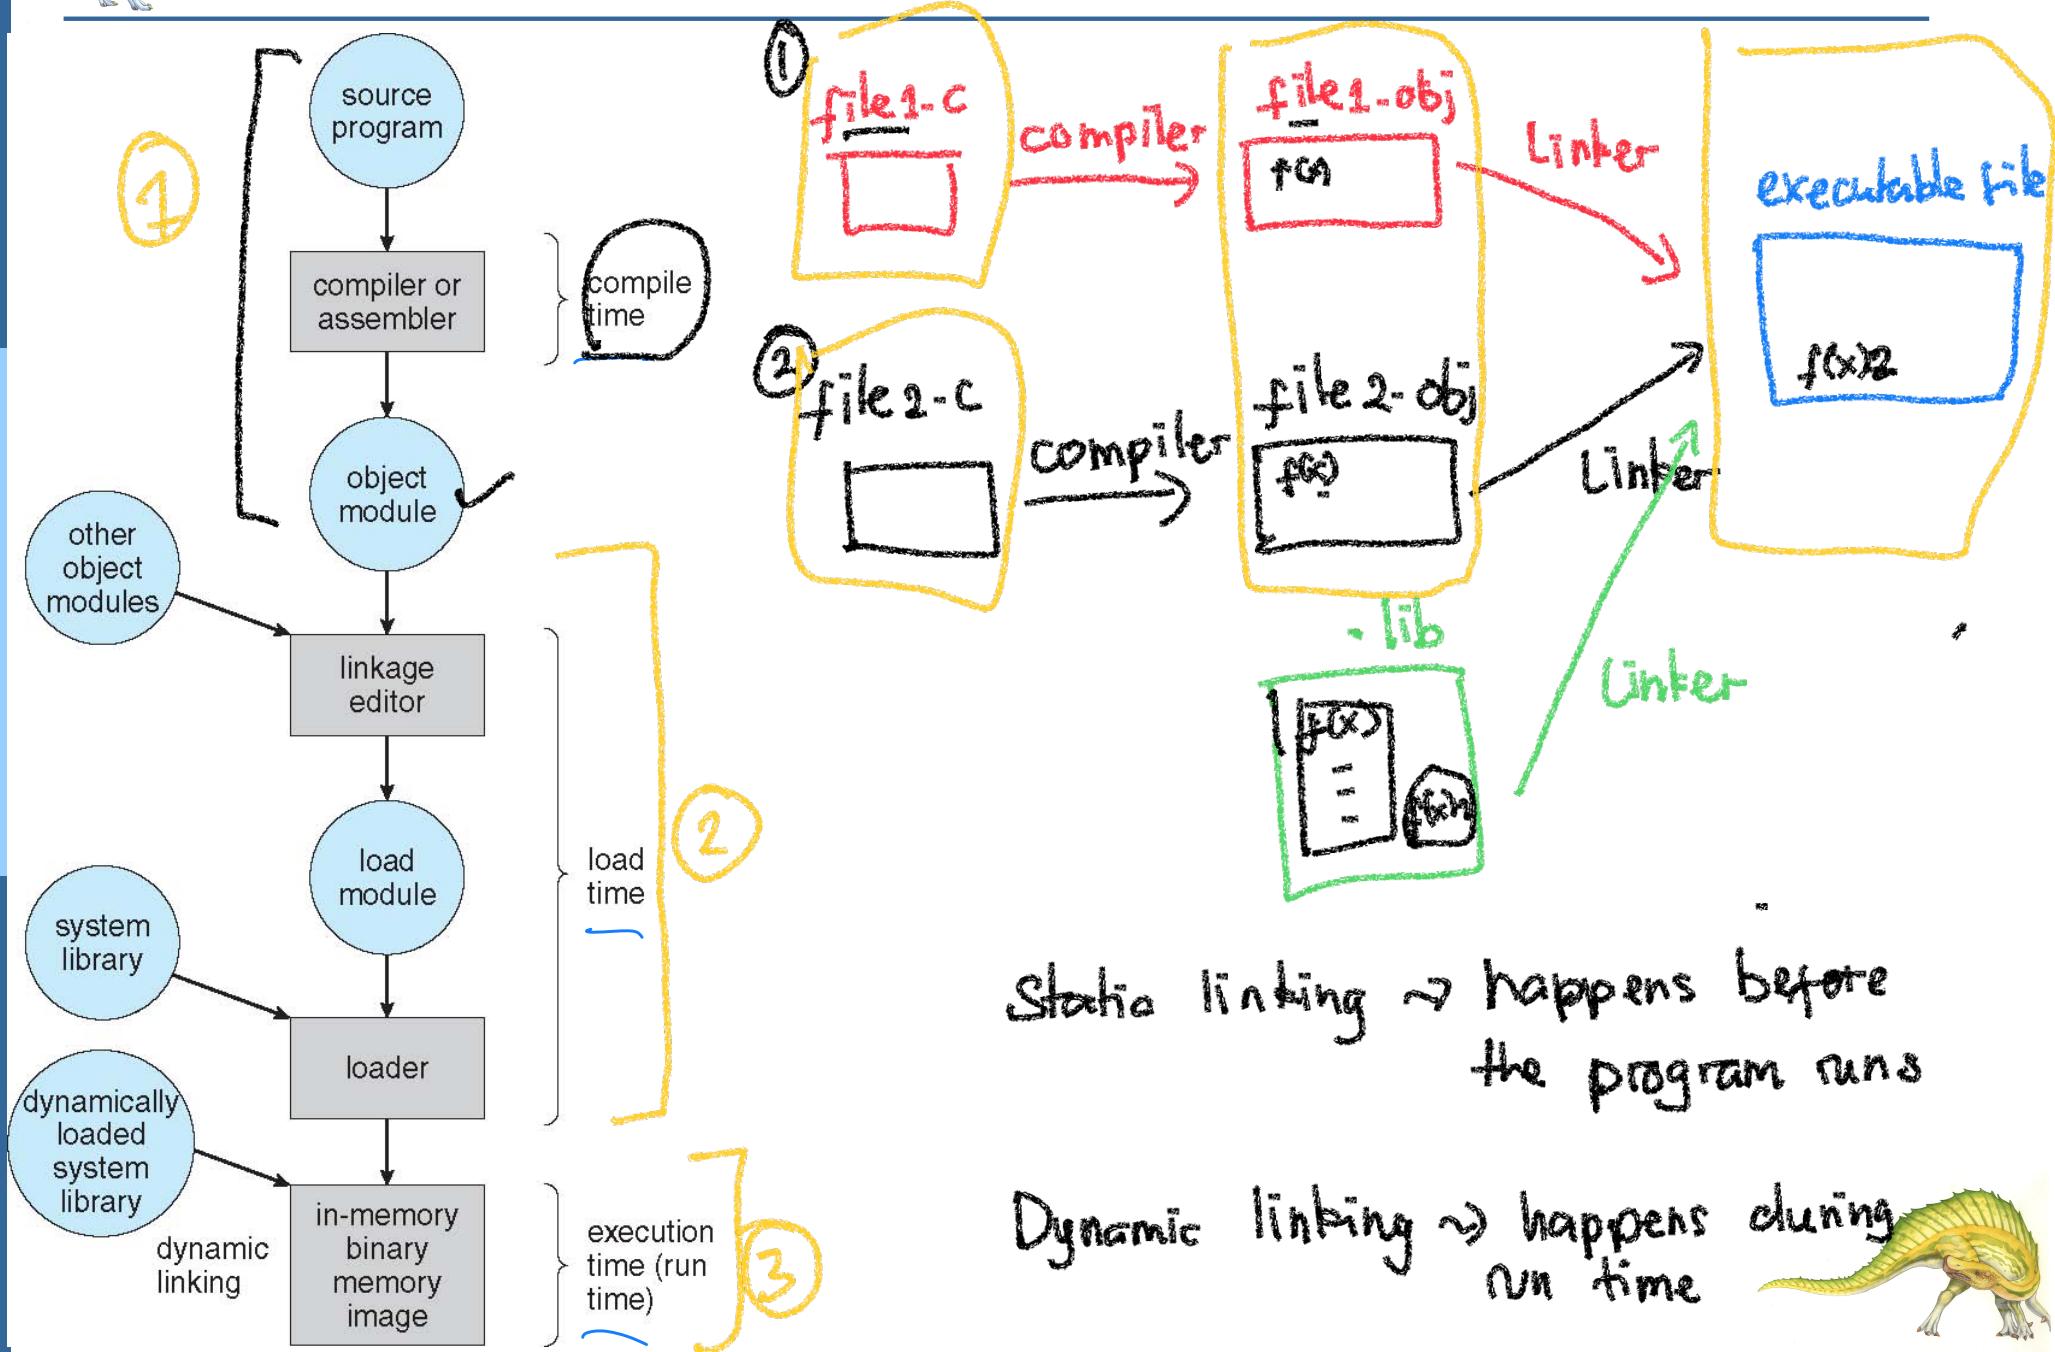
Binding of Instructions and Data to Memory

- Address binding of instructions and data to memory addresses can happen at three different stages
 - **Compile time:** If memory location known a priori, **absolute code** can be generated; must recompile code if starting location changes
 - **Load time:** Must generate **relocatable code** if memory location is not known at compile time
 - **Execution time:** Binding delayed until run time if the process can be moved during its execution from one memory segment to another
 - Need hardware support for address maps (e.g., base and limit registers)





Multistep Processing of a User Program





Logical vs. Physical Address Space

- The concept of a logical address space that is bound to a separate **physical address space** is central to proper memory management
 - **Logical address** – generated by the CPU; also referred to as **virtual address**
 - **Physical address** – address seen by the memory unit
- Logical and physical addresses are the same in compile-time and load-time address-binding schemes; logical (virtual) and physical addresses differ in execution-time address-binding scheme
- **Logical address space** is the set of all logical addresses generated by a program
- **Physical address space** is the set of all physical addresses generated by a program





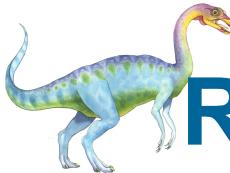
Logical vs. Physical Address Space

- The user program deals with *logical addresses*; it never sees the *real* physical addresses
 - Execution-time binding occurs when reference is made to location in memory
 - Logical address bound to physical addresses

```
#include <stdio.h>
#include <stdlib.h>

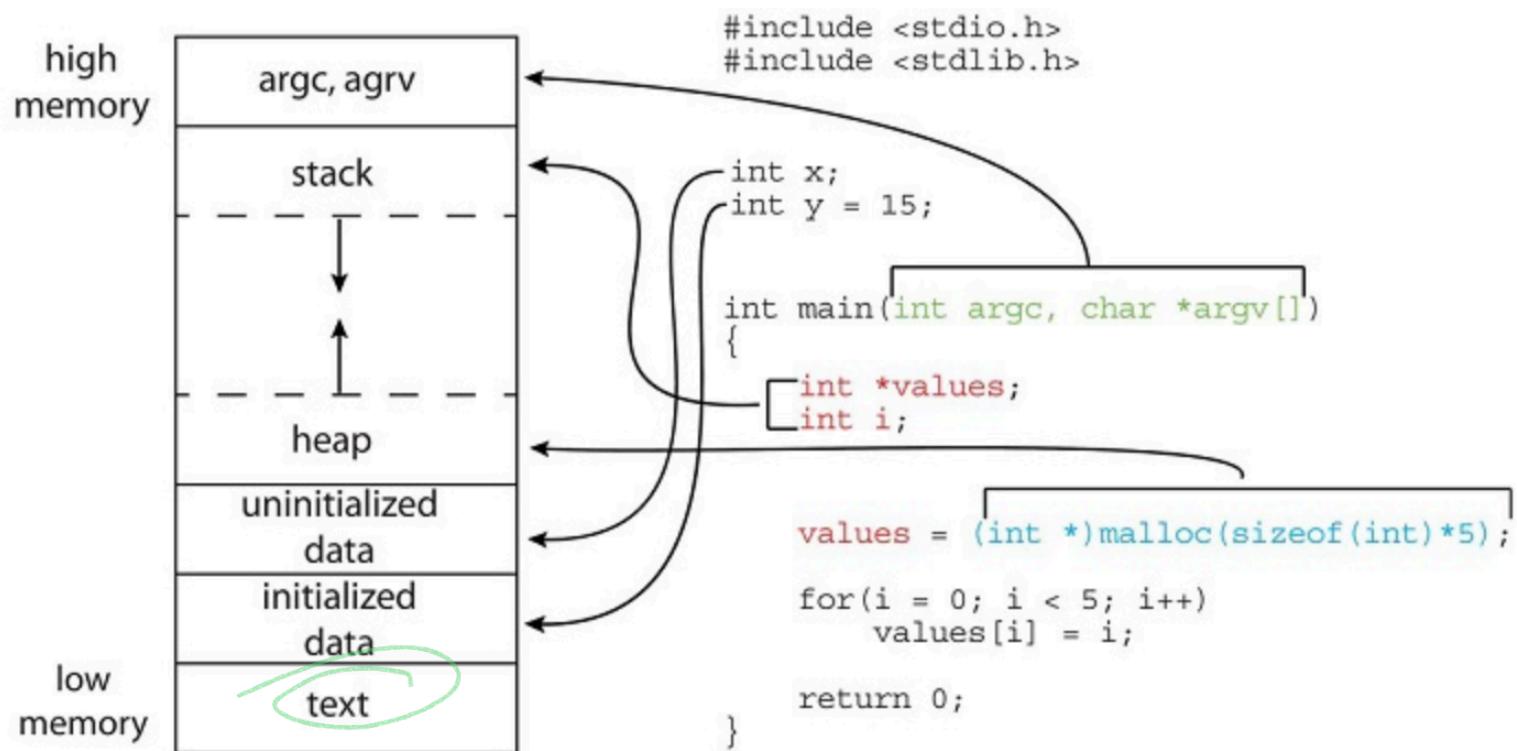
int main(int argc, char *argv[]) {
    printf("location of code : %p\n", main);
    printf("location of heap : %p\n", malloc(100e6));
    int x = 3;
    printf("location of stack: %p\n", &x);
    return 0;
}
```





Recall Memory Layout of a C Program

- The global data section is divided into different sections for (a) initialized data and (b) uninitialized data.
- A separate section is provided for the `argc` and `argv` parameters passed to the `main()` function.





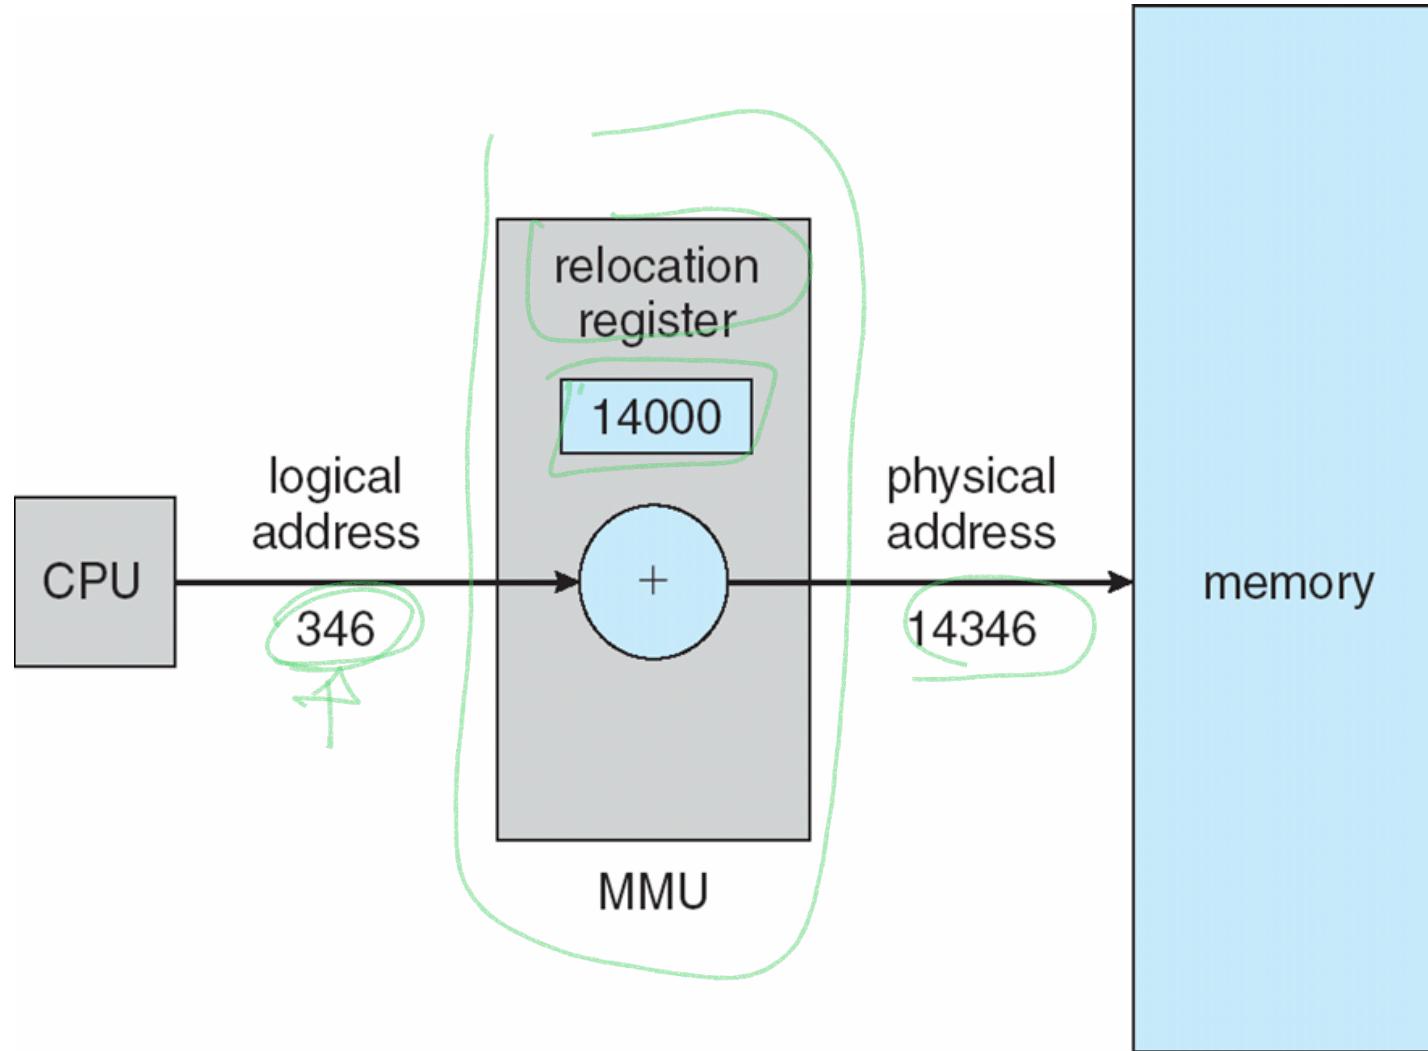
Memory-Management Unit (MMU)

- Hardware device that at run time maps virtual to physical address
- Many methods possible, covered in the rest of this chapter
- To start, consider simple scheme where the value in the relocation register is added to every address generated by a user process at the time it is sent to memory
 - Base register now called **relocation register**
 - MS-DOS on Intel 80x86 used 4 relocation registers





MMU





Dynamic relocation using a relocation register

- Routine is not loaded until it is called
- Better memory-space utilization; unused routine is never loaded
- All routines kept on disk in relocatable load format
- Useful when large amounts of code are needed to handle infrequently occurring cases
- No special support from the operating system is required
 - Implemented through program design
 - OS can help by providing libraries to implement dynamic loading

-dll ~> Unix } for dynamic library
 .so ~> windows }





Dynamic Linking

- **Static linking** – system libraries and program code combined by the loader into the binary program image
- Dynamic linking –linking postponed until execution time
- Small piece of code, **stub**, used to locate the appropriate memory-resident library routine
- Stub replaces itself with the address of the routine, and executes the routine
- Operating system checks if routine is in processes' memory address
 - If not in address space, add to address space
- Dynamic linking is particularly useful for libraries
- System also known as **shared libraries**
- Consider applicability to patching system libraries
 - Versioning may be needed





Swapping

- A process can be **swapped** temporarily out of memory to a backing store, and then brought back into memory for continued execution
 - Total physical memory space of processes can exceed physical memory
- **Backing store** – fast disk large enough to accommodate copies of all memory images for all users; must provide direct access to these memory images
- **Roll out, roll in** – swapping variant used for priority-based scheduling algorithms; lower-priority process is swapped out so higher-priority process can be loaded and executed
- Major part of swap time is transfer time; total transfer time is directly proportional to the amount of memory swapped
- System maintains a **ready queue** of ready-to-run processes which have memory images on disk





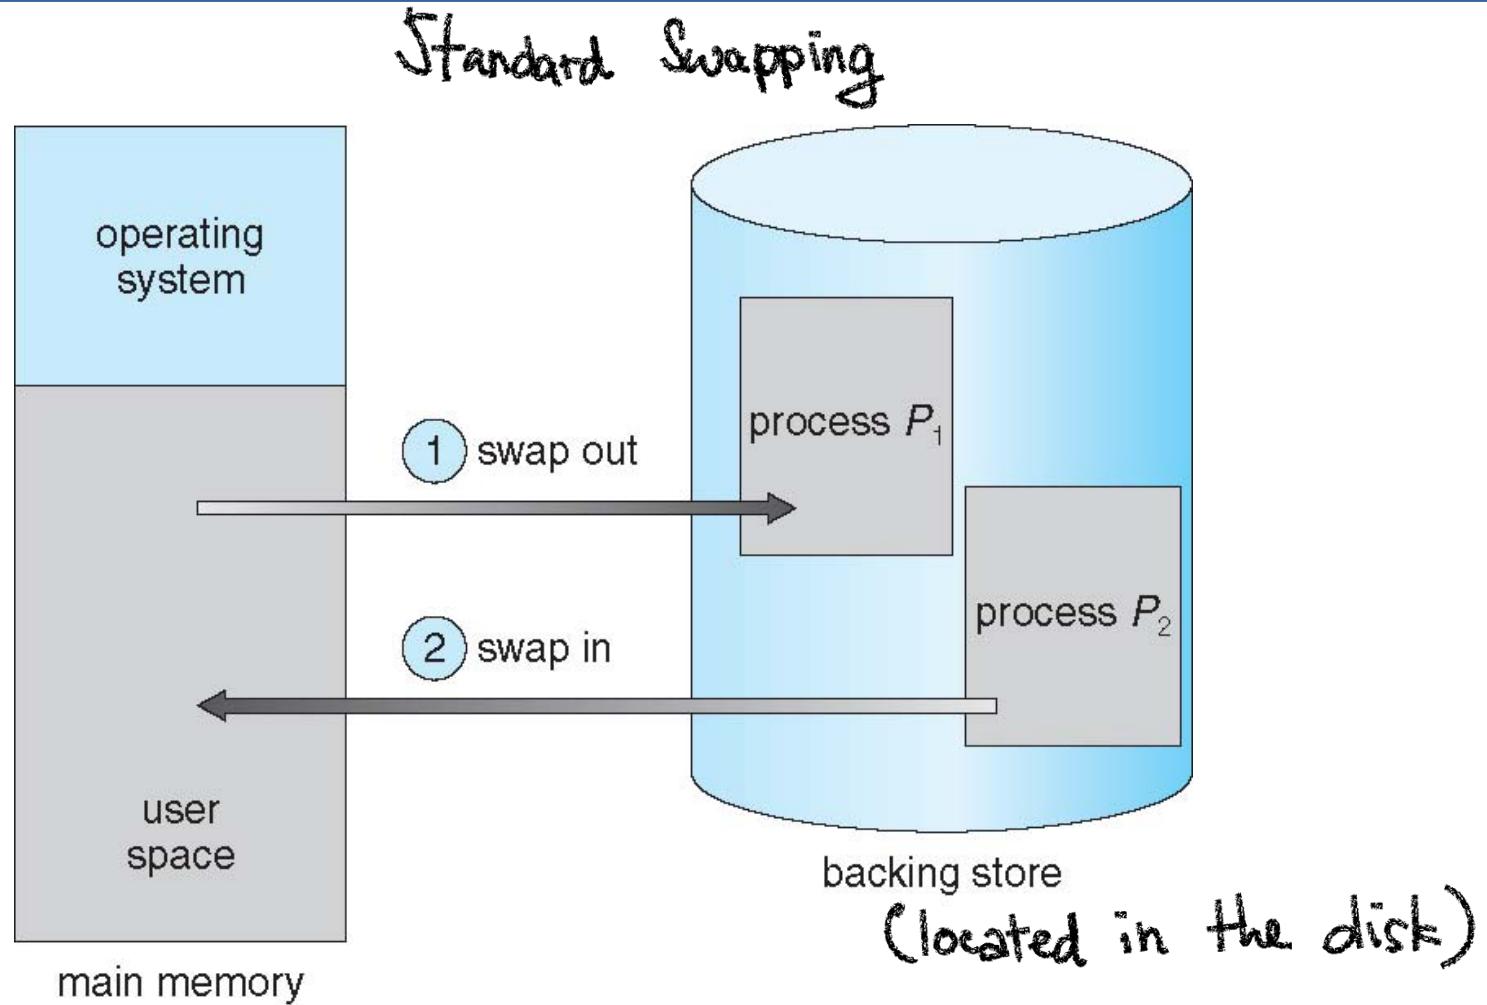
Swapping (Cont.)

- Does the swapped out process need to swap back in to same physical addresses?
- Depends on address binding method
 - Plus consider pending I/O to / from process memory space
- Modified versions of swapping are found on many systems (i.e., UNIX, Linux, and Windows)
 - Swapping normally disabled
 - Started if more than threshold amount of memory allocated
 - Disabled again once memory demand reduced below threshold





Schematic View of Swapping



Swapping is expensive , why ? disk is slow





Context Switch Time including Swapping

- If next processes to be put on CPU is not in memory, need to swap out a process and swap in target process
- Context switch time can then be very high
- 100MB process swapping to hard disk with transfer rate of 50MB/sec
 - Swap out time of 2000 ms
 - Plus swap in of same sized process
 - Total context switch swapping component time of 4000ms (4 seconds)
- Can reduce if reduce size of memory swapped – by knowing how much memory really being used
 - System calls to inform OS of memory use via `request_memory()` and `release_memory()`

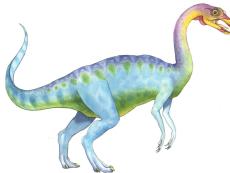




Context Switch Time and Swapping (Cont.)

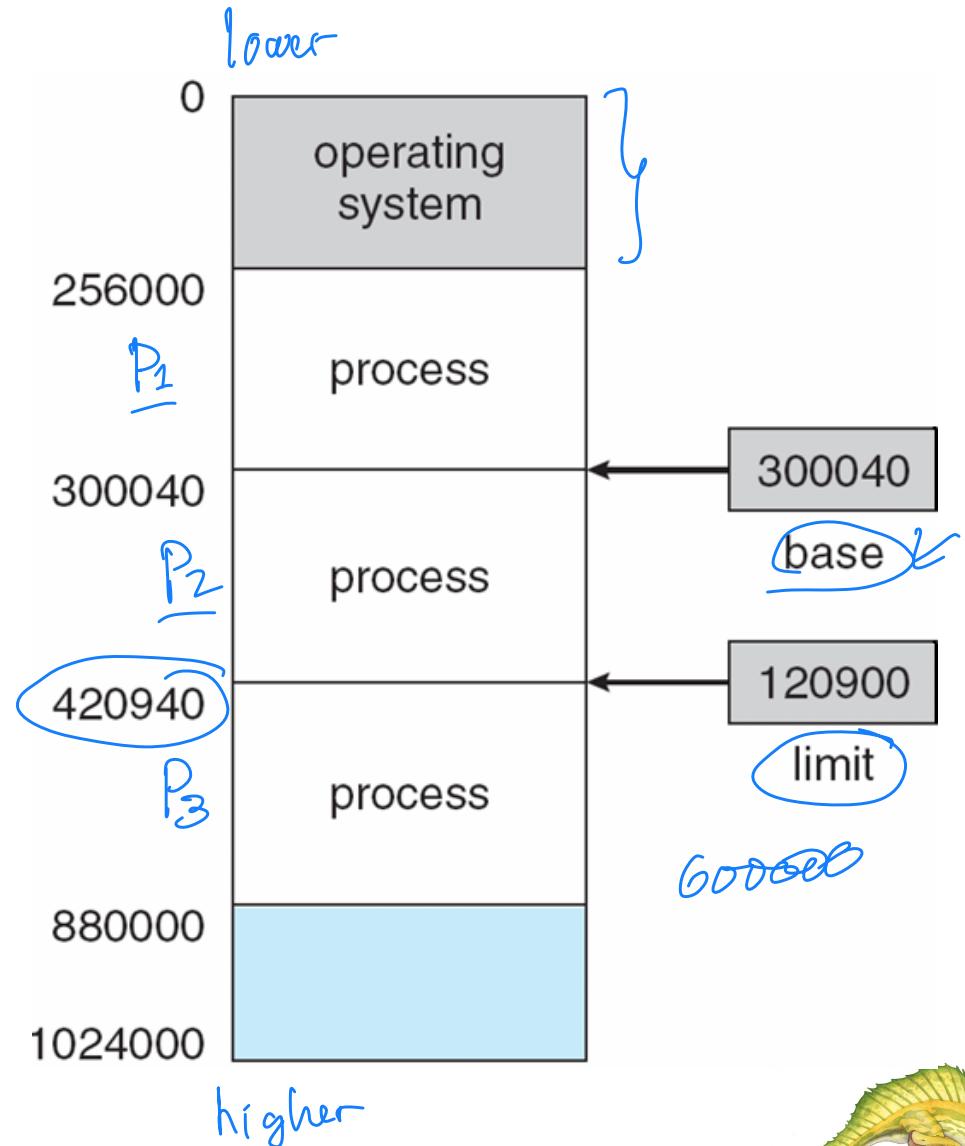
- Other constraints as well on swapping
 - Pending I/O – can't swap out as I/O would occur to wrong process
 - Or always transfer I/O to kernel space, then to I/O device
 - Known as **double buffering**, adds overhead
- Standard swapping not used in modern operating systems
 - But modified version common
 - **Swap only when free memory extremely low**





Contiguous Allocation

- Main memory must support both OS and user processes
- Limited resource, must allocate efficiently
- Contiguous allocation is one early method \leadsto not efficient
- Main memory usually into two **partitions**:
 - Resident operating system, usually held in low memory with interrupt vector
 - User processes then held in high memory
 - Each process contained in single contiguous section of memory

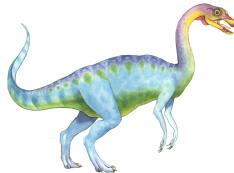




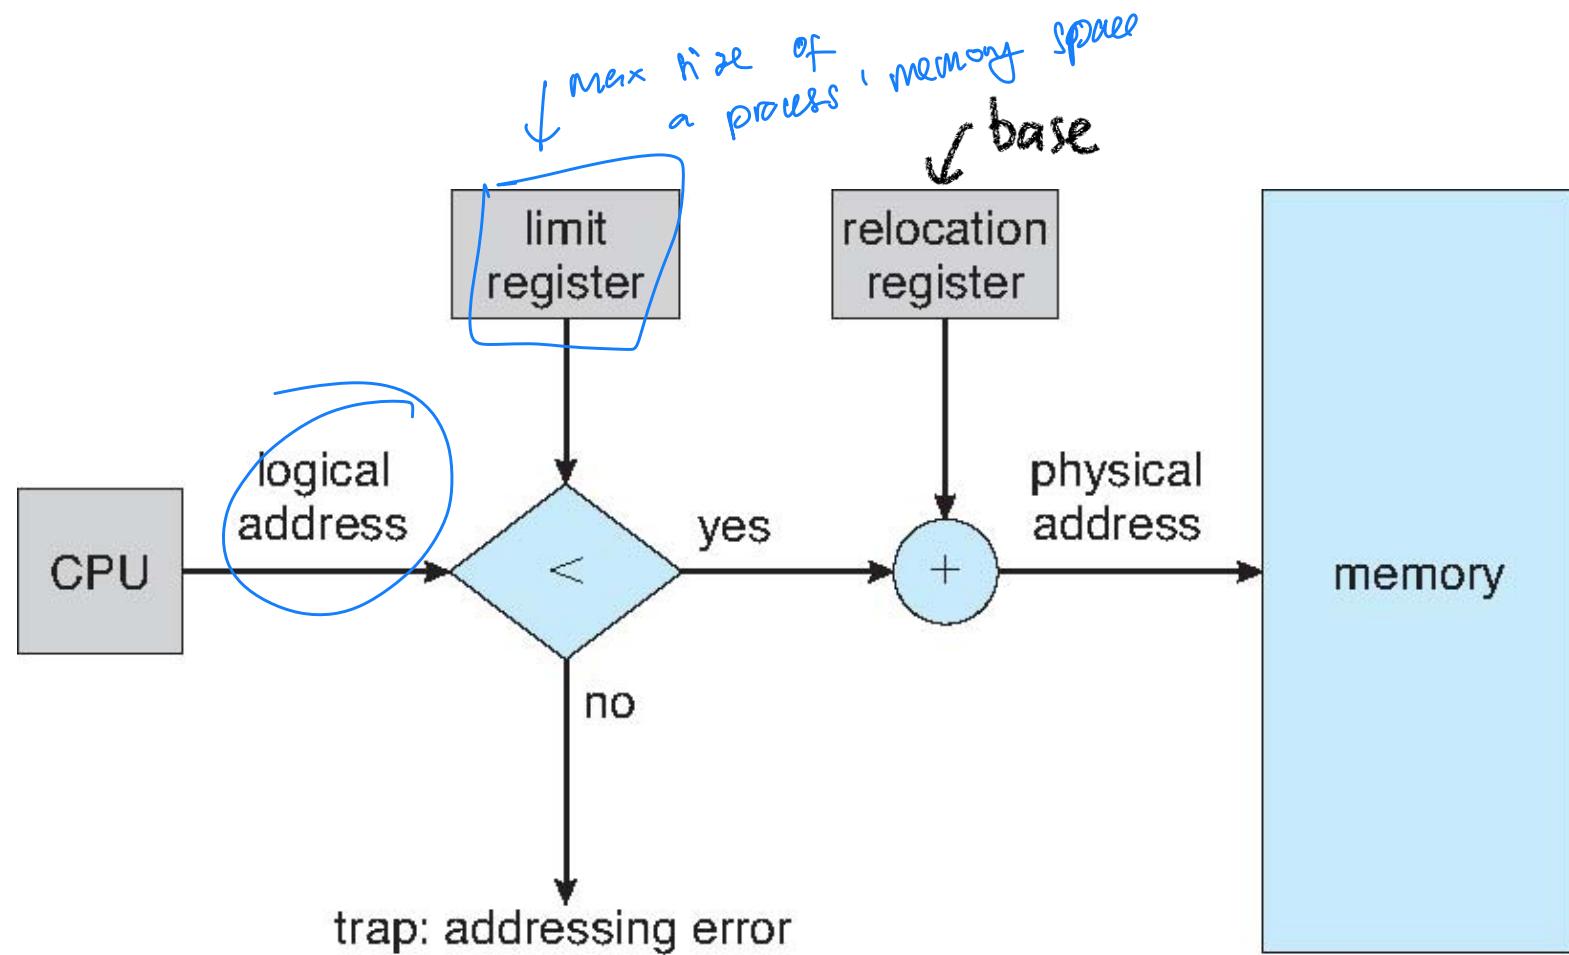
Contiguous Allocation (Cont.)

- Relocation registers used to protect user processes from each other, and from changing operating-system code and data
 - Base register contains value of smallest physical address
 - Limit register contains range of logical addresses – each logical address must be less than the limit register
 - MMU maps logical address *dynamically*
 - Can then allow actions such as kernel code being **transient** and kernel changing size





Hardware Support for Relocation and Limit Registers

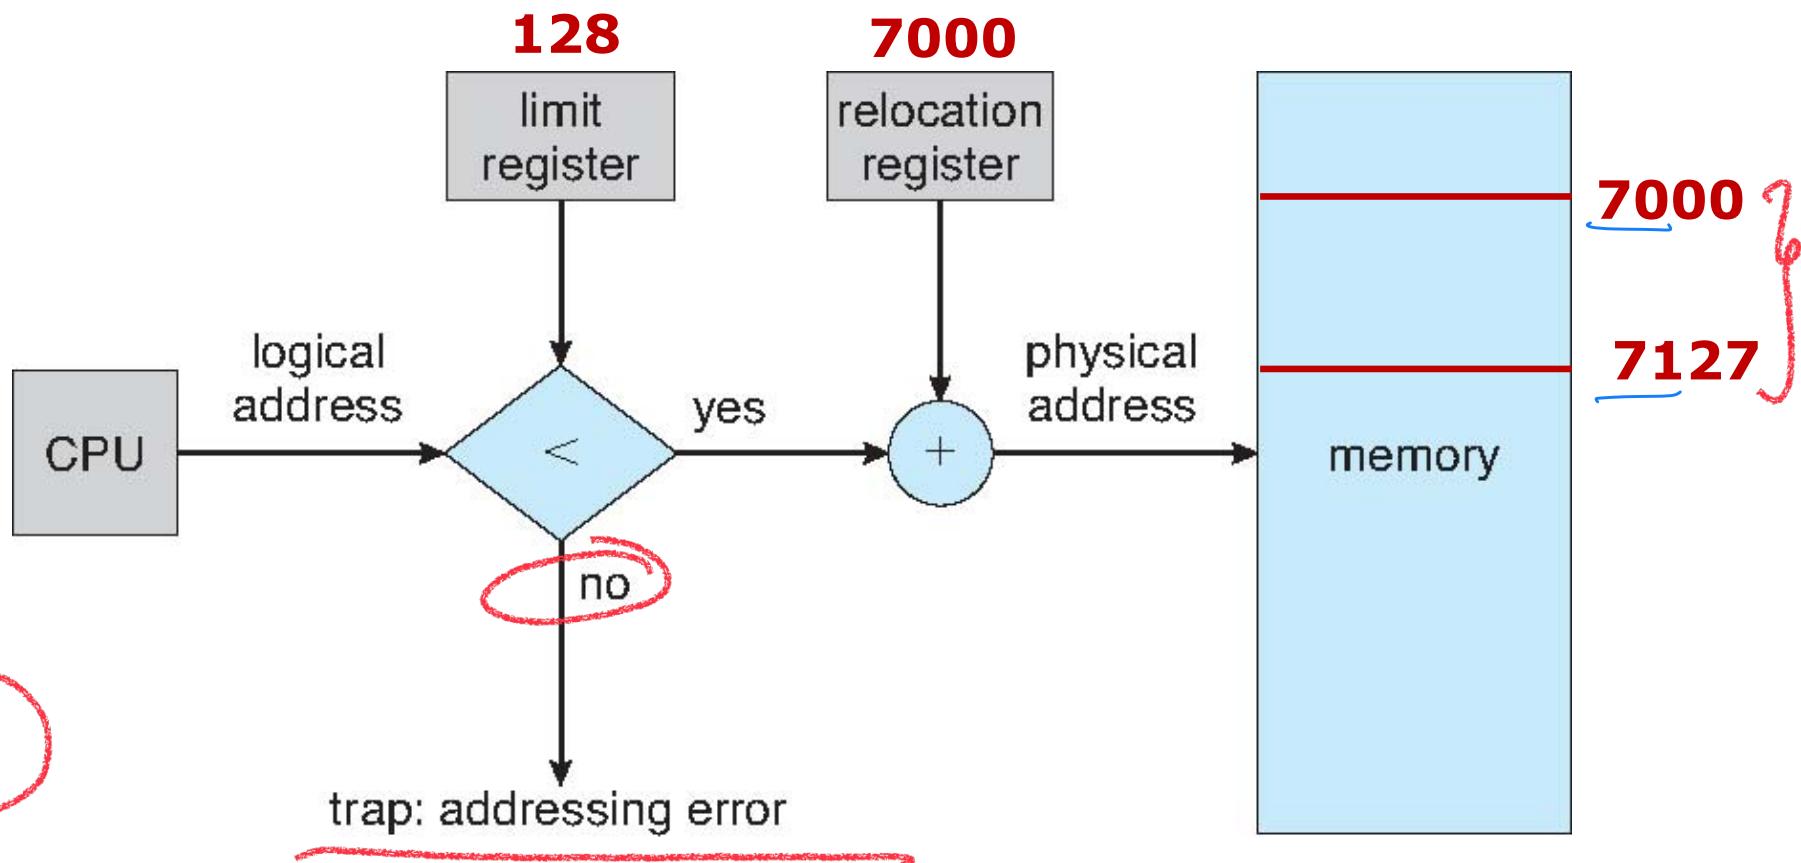


P₃:





Hardware Support for Relocation and Limit Registers



Suppose a process requires max 128 bytes and relocation register is set = 7000

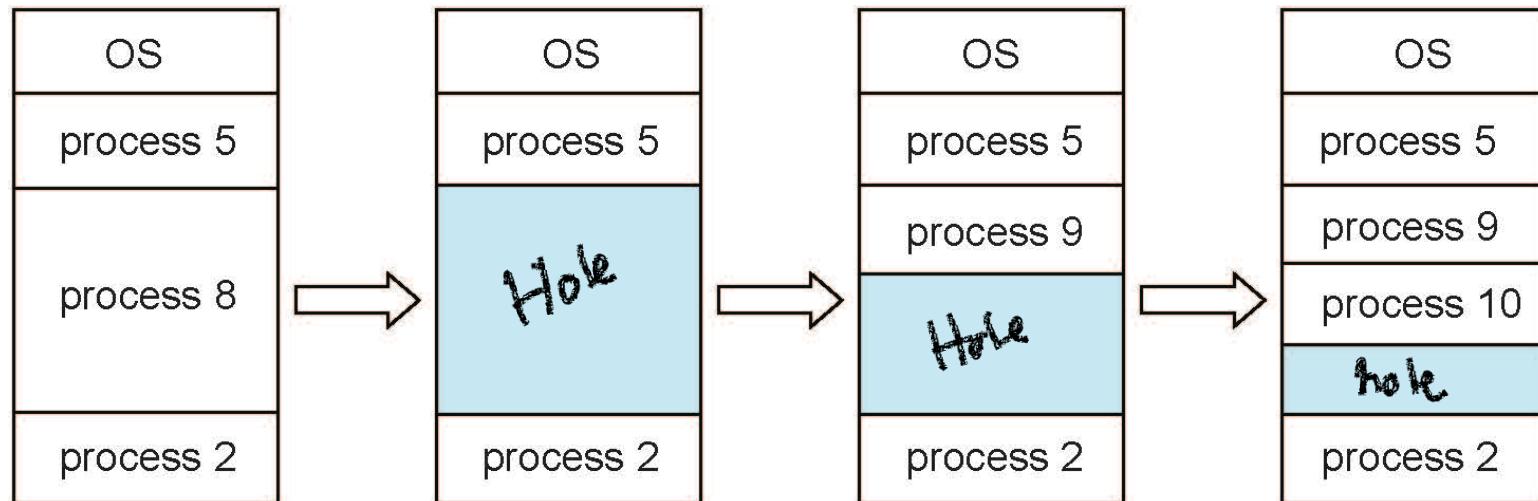
In most modern systems, memory is **byte-addressable**, meaning each memory address points to a single byte (8 bits).





Multiple-partition allocation

- Multiple-partition allocation
 - Degree of multiprogramming limited by number of partitions
 - **Variable-partition** sizes for efficiency (sized to a given process' needs)
 - **Hole** – block of available memory; holes of various size are scattered throughout memory
 - When a process arrives, it is allocated memory from a hole large enough to accommodate it
 - Process exiting frees its partition, adjacent free partitions combined
 - Operating system maintains information about:
 - a) allocated partitions
 - b) free partitions (hole)





Dynamic Storage-Allocation Problem

How to satisfy a request of size n from a list of free holes?

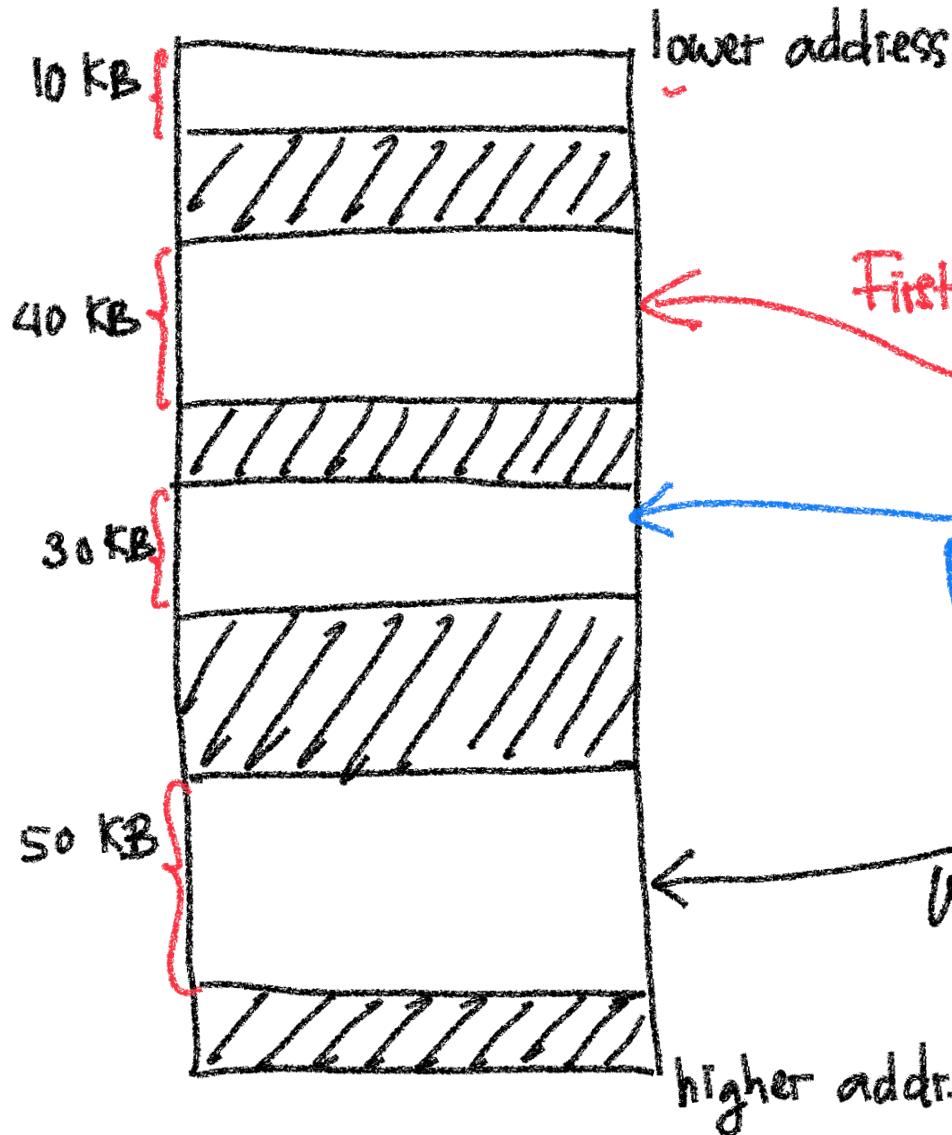
- **First-fit:** Allocate the *first* hole that is big enough
- **Best-fit:** Allocate the *smallest* hole that is big enough; must search entire list, unless ordered by size
 - Produces the smallest leftover hole
- **Worst-fit:** Allocate the *largest* hole; must also search entire list
 - Produces the largest leftover hole

First-fit and best-fit better than worst-fit in terms of speed and storage utilization





Memory



First Fit

Suppose a process needs

20 KB

Best Fit

Worst Fit





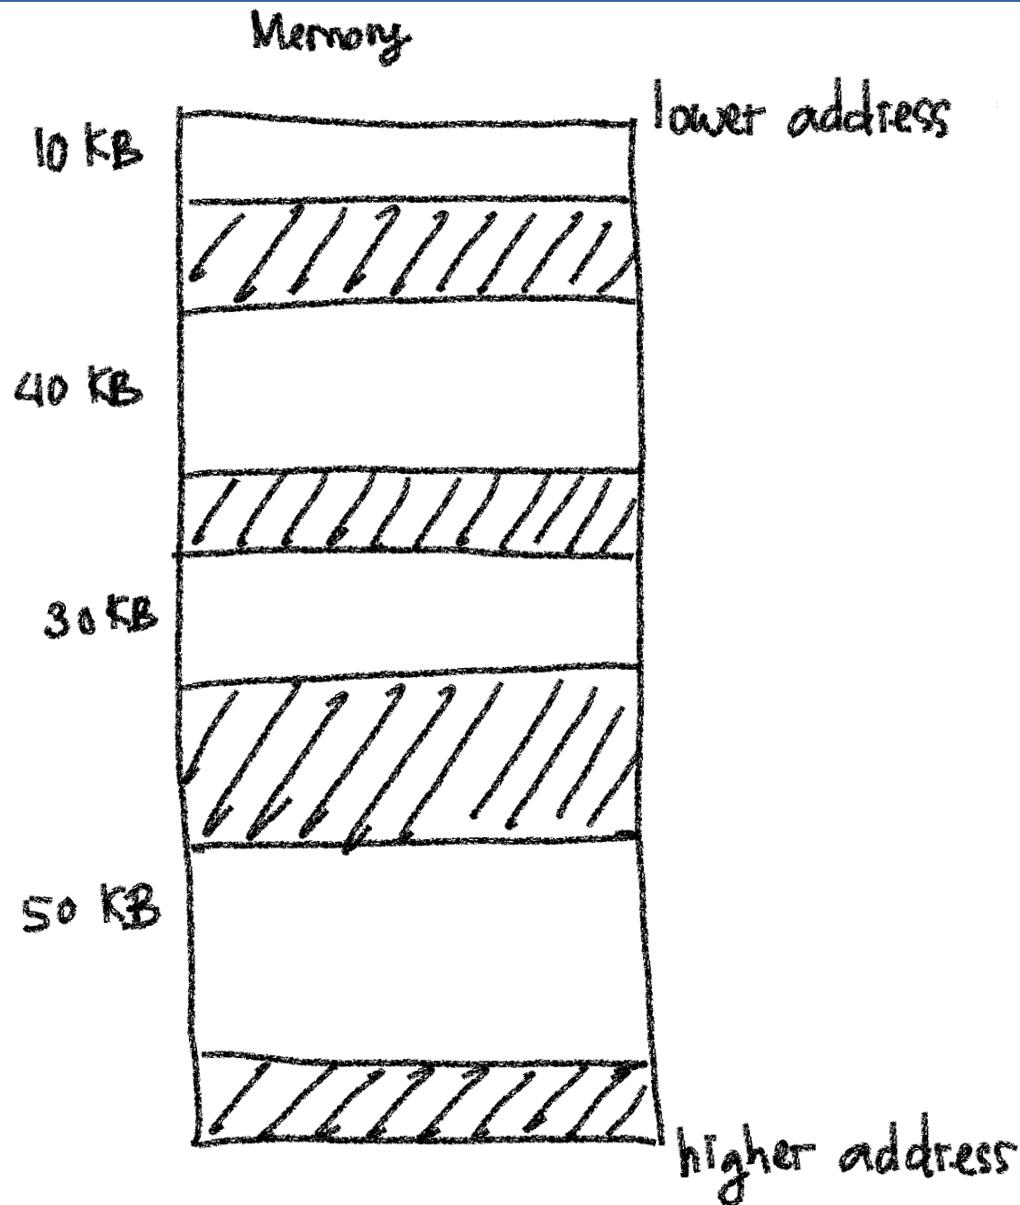
Fragmentation

- **External Fragmentation** – total memory space exists to satisfy a request, but it is not contiguous
- First fit analysis reveals that given N blocks allocated, $0.5 N$ blocks lost to fragmentation
 - 1/3 may be unusable -> **50-percent rule**





External Fragmentation



Suppose a process needs

130 KB





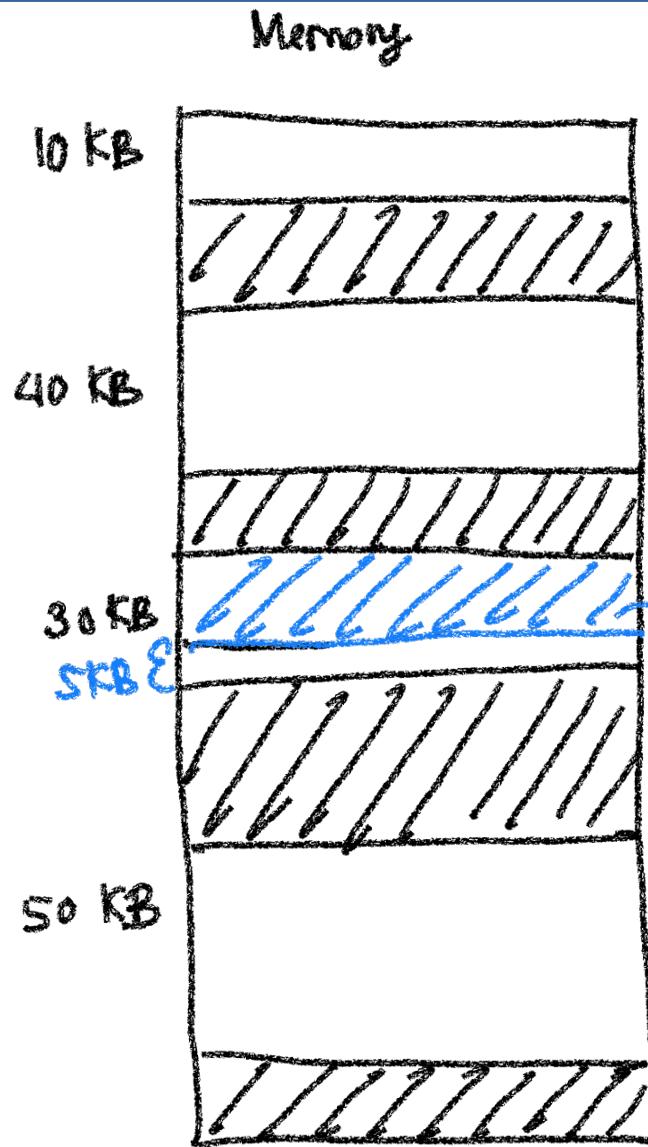
Fragmentation

- **Internal Fragmentation** – allocated memory may be slightly larger than requested memory; this size difference is memory internal to a partition, but not being used
- Consider a multiple-partition allocation scheme with a hole of 18,464 bytes. Suppose that the next process requests 18,462 bytes. If we allocate exactly the requested block, we are left with a hole of 2 bytes. The overhead to keep track of this hole will be substantially larger than the hole itself.
- The general approach to avoiding this problem is to break the physical memory into fixed-sized blocks and allocate memory in units based on block size.





Internal Fragmentation



lower address

Suppose a process needs

25 KB

Best Fit





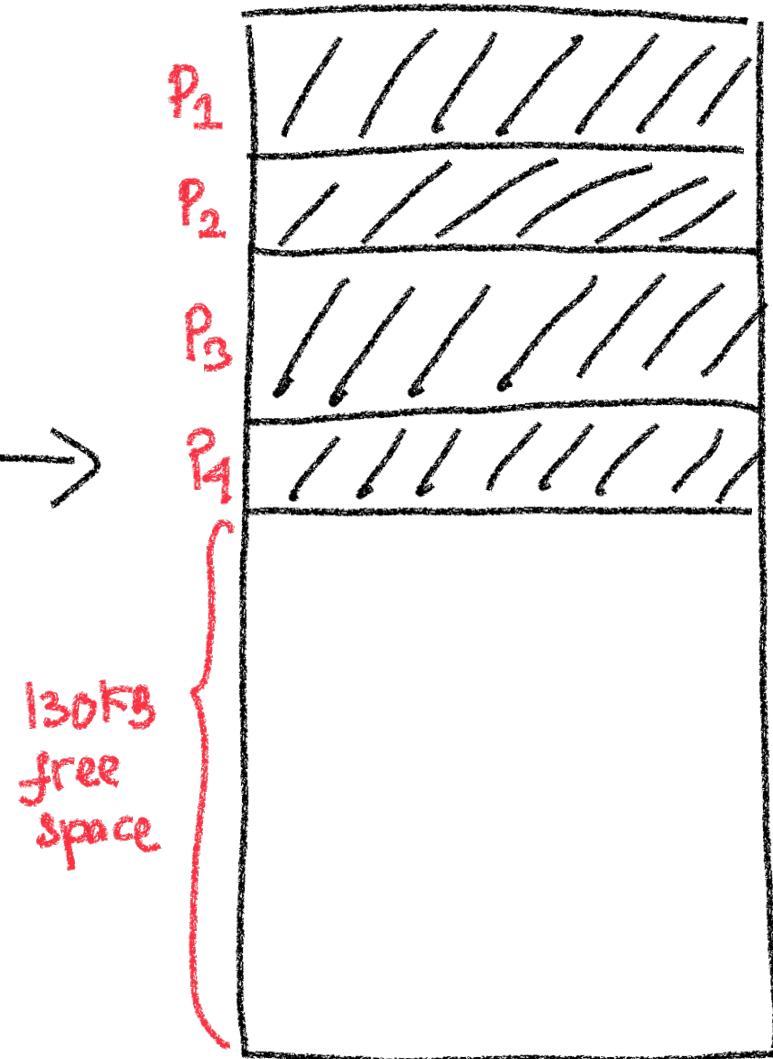
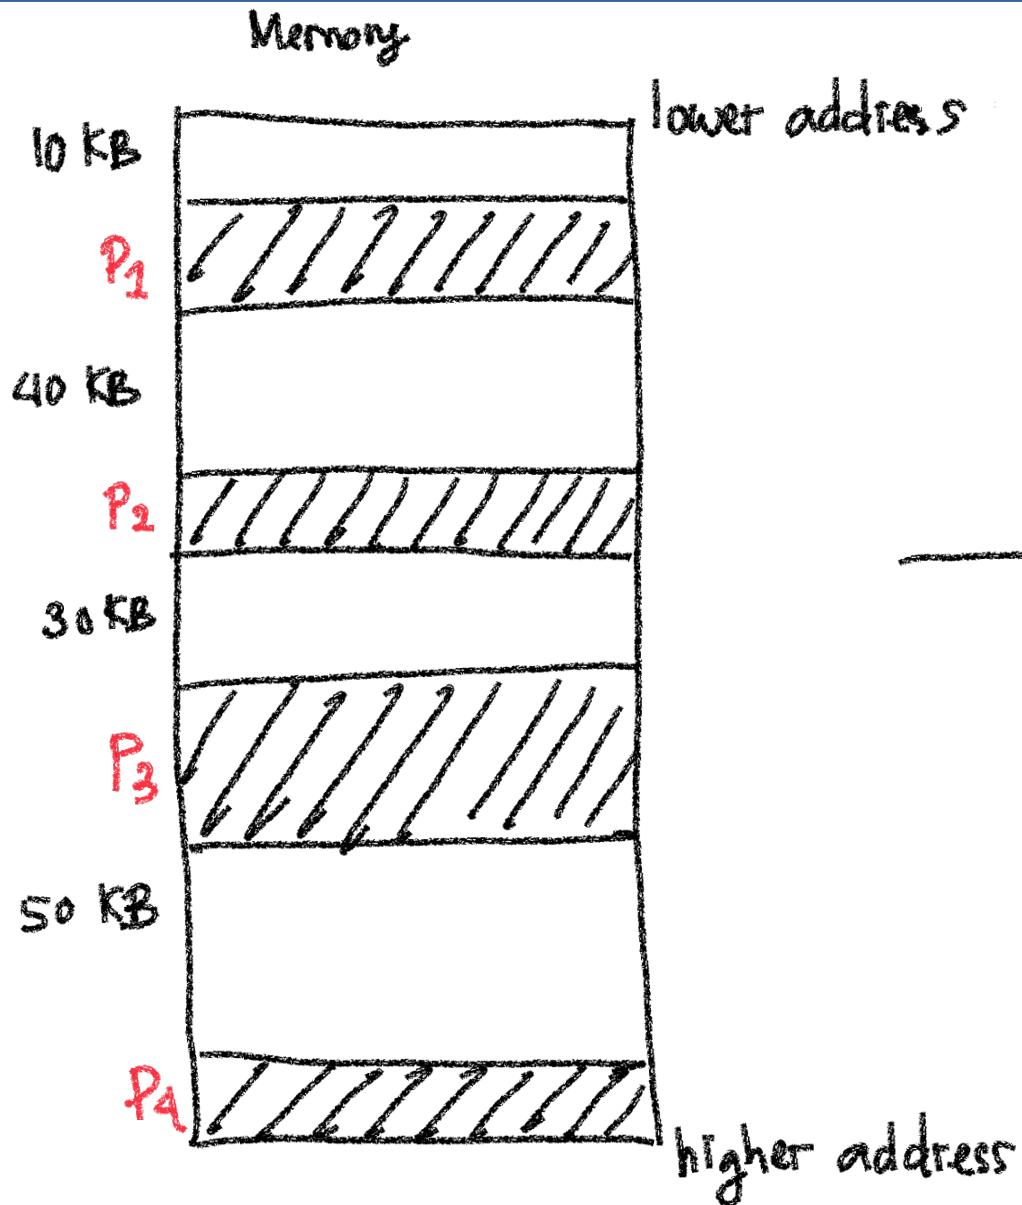
Fragmentation (Cont.)

- Reduce external fragmentation by **compaction**
 - Shuffle memory contents to place all free memory together in one large block
 - Compaction is possible *only* if relocation is dynamic, and is done at execution time
- Now consider that backing store has same fragmentation problems

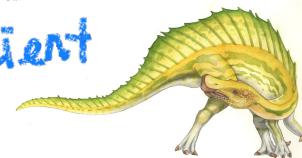




Compaction



Still not efficient
why?





Paging

8 bytes:
~~3~~

2 bytes

$$2^0 = 1$$

$$2^1 = 2$$

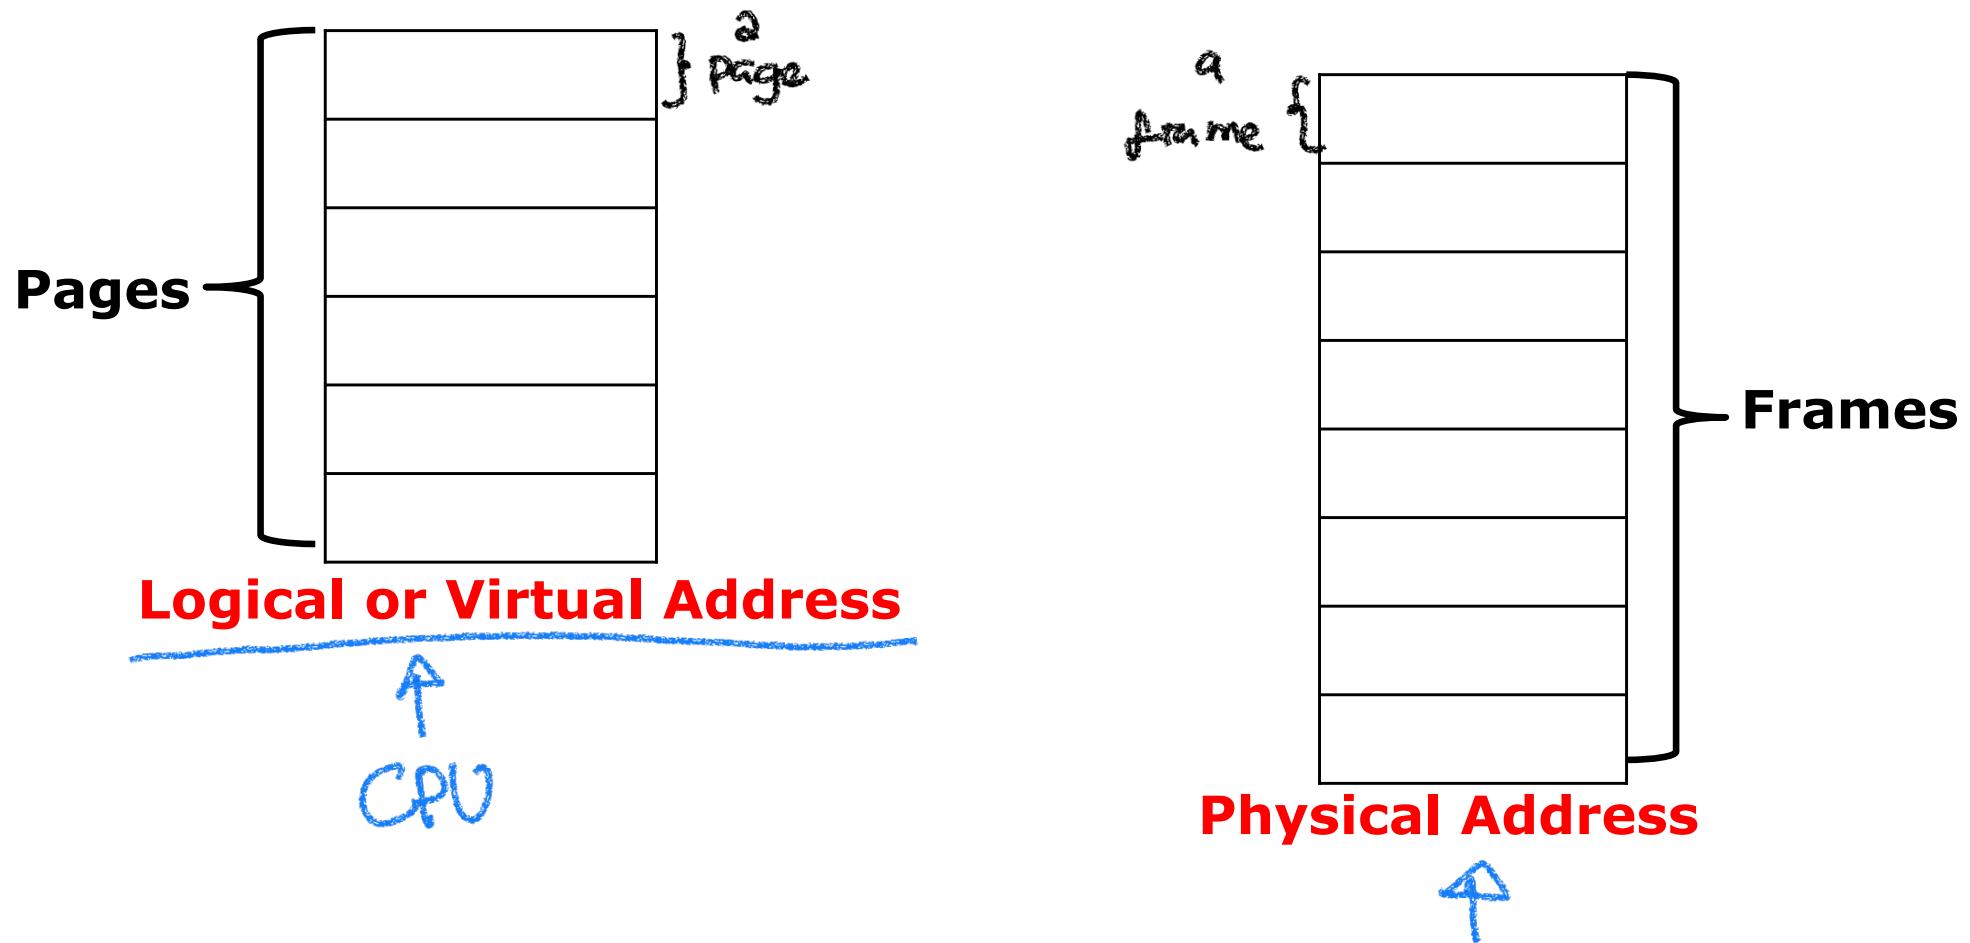
$$2^2 = 4$$

$$2^3 = 8$$

- Physical address space of a process can be noncontiguous; process is allocated physical memory whenever the latter is available
 - Avoids external fragmentation
 - Avoids problem of varying sized memory chunks
- Divide physical memory into fixed-sized blocks called **frames**
 - Size is power of 2, between 512 bytes and 16 Mbytes \rightarrow
- Divide logical memory into blocks of same size called **pages**
- Keep track of all free frames
- To run a program of size **N** pages, need to find **N** free frames and load program
- Set up a **page table** (per process) to translate logical to physical addresses
- Backing store (**used for Swapping**) likewise split into pages
- Still have Internal fragmentation

Lo

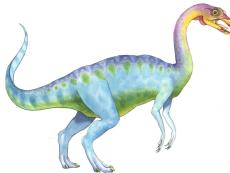




page size = frame size

a page → a frame





Aside:

1 KB

A frame size is power of 2, between **512 bytes** (**2⁹ bytes**) and **16 Mbytes** ($2^4 * 2^{20} = 2^{4+20} = 2^{24}$ bytes)

Computers use binary addressing, meaning that each memory address corresponds to a unique binary number.

e.g. 6 bits space:

1	0	0	1	0	0
---	---	---	---	---	---

2^6 Bytes Possible addresses

$$\underline{1 \text{ KB} = 2^{10} \text{ Bytes}}$$

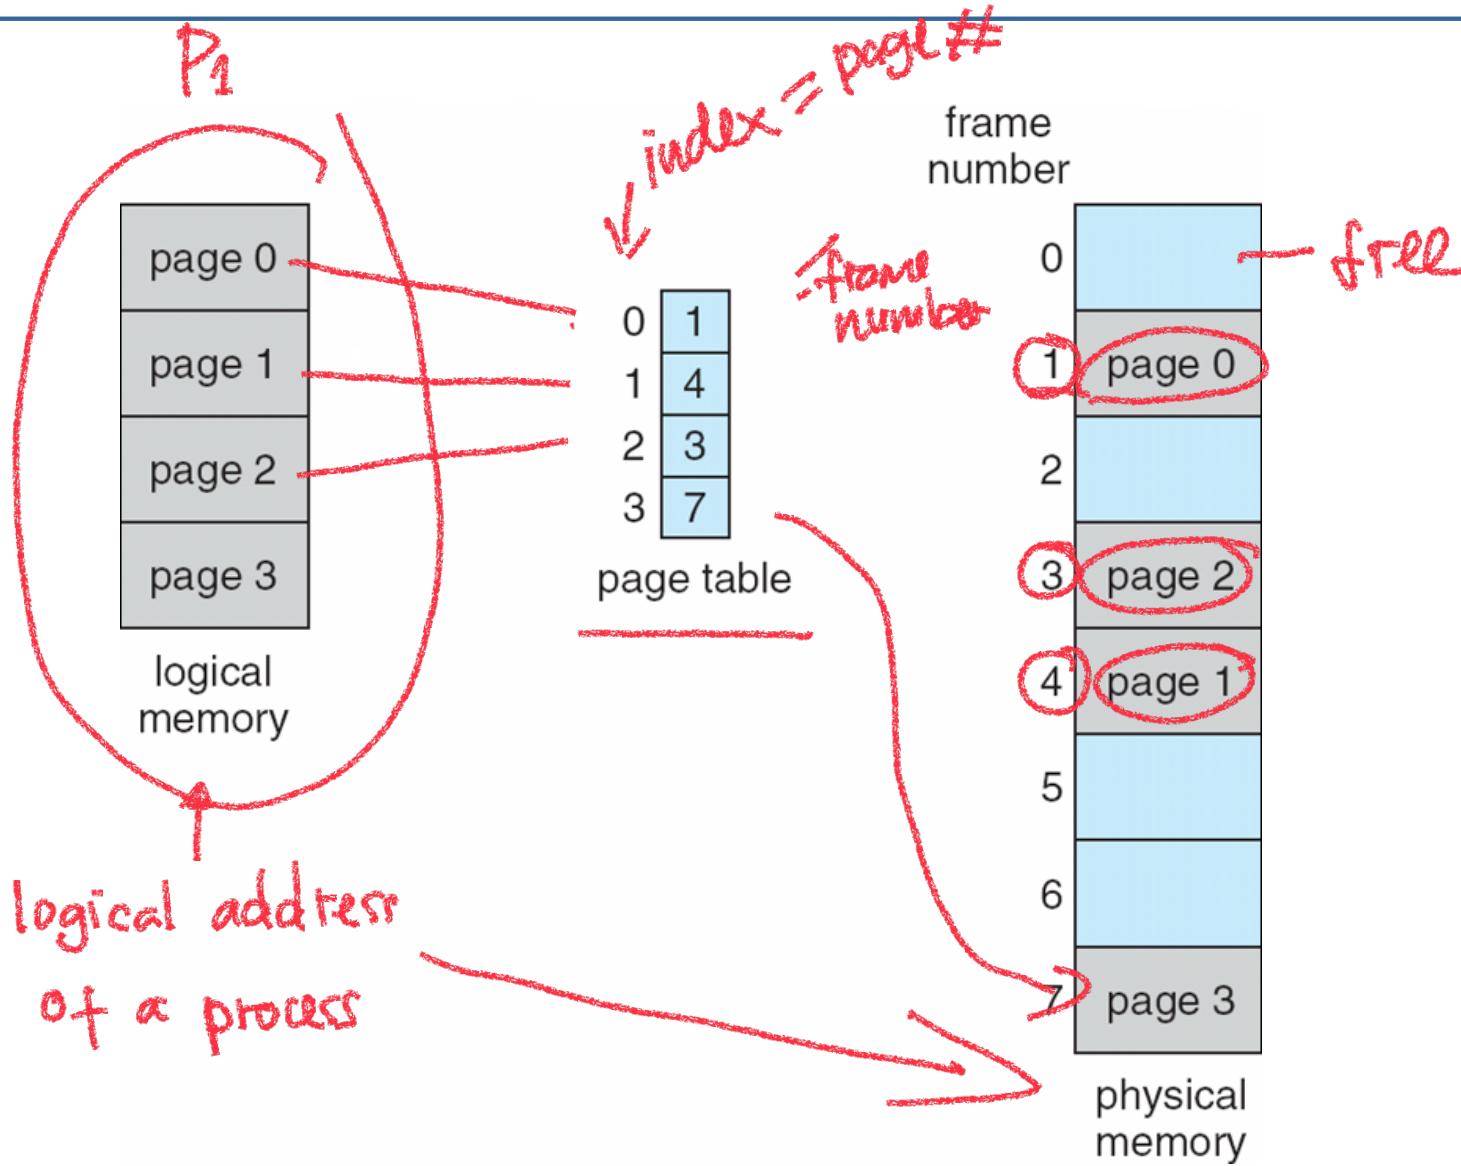
$$1 \text{ MB} = 2^{20} \text{ Bytes}$$

$$1 \text{ GB} = 2^{30} \text{ Bytes}$$





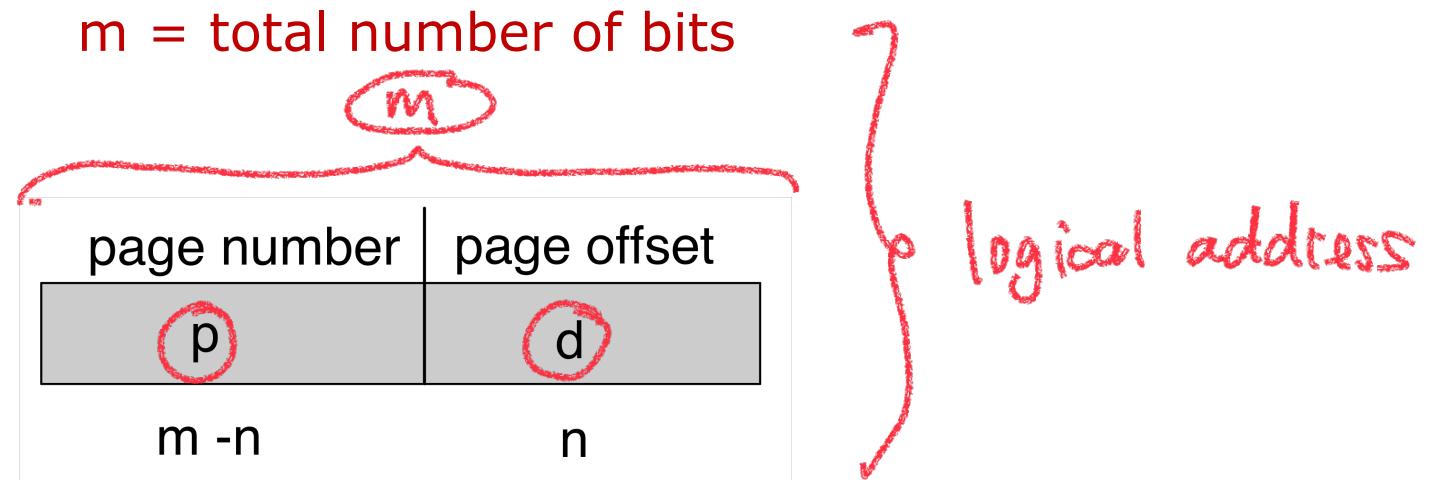
Paging Model of Logical and Physical Memory





Address Translation Scheme

- Address generated by CPU (**logical address**) is divided into:
 - **Page number (p)** – used as an index into a **page table** which contains base address of each page in physical memory
 - **Page offset (d)** – combined with base address to define the physical memory address that is sent to the memory unit

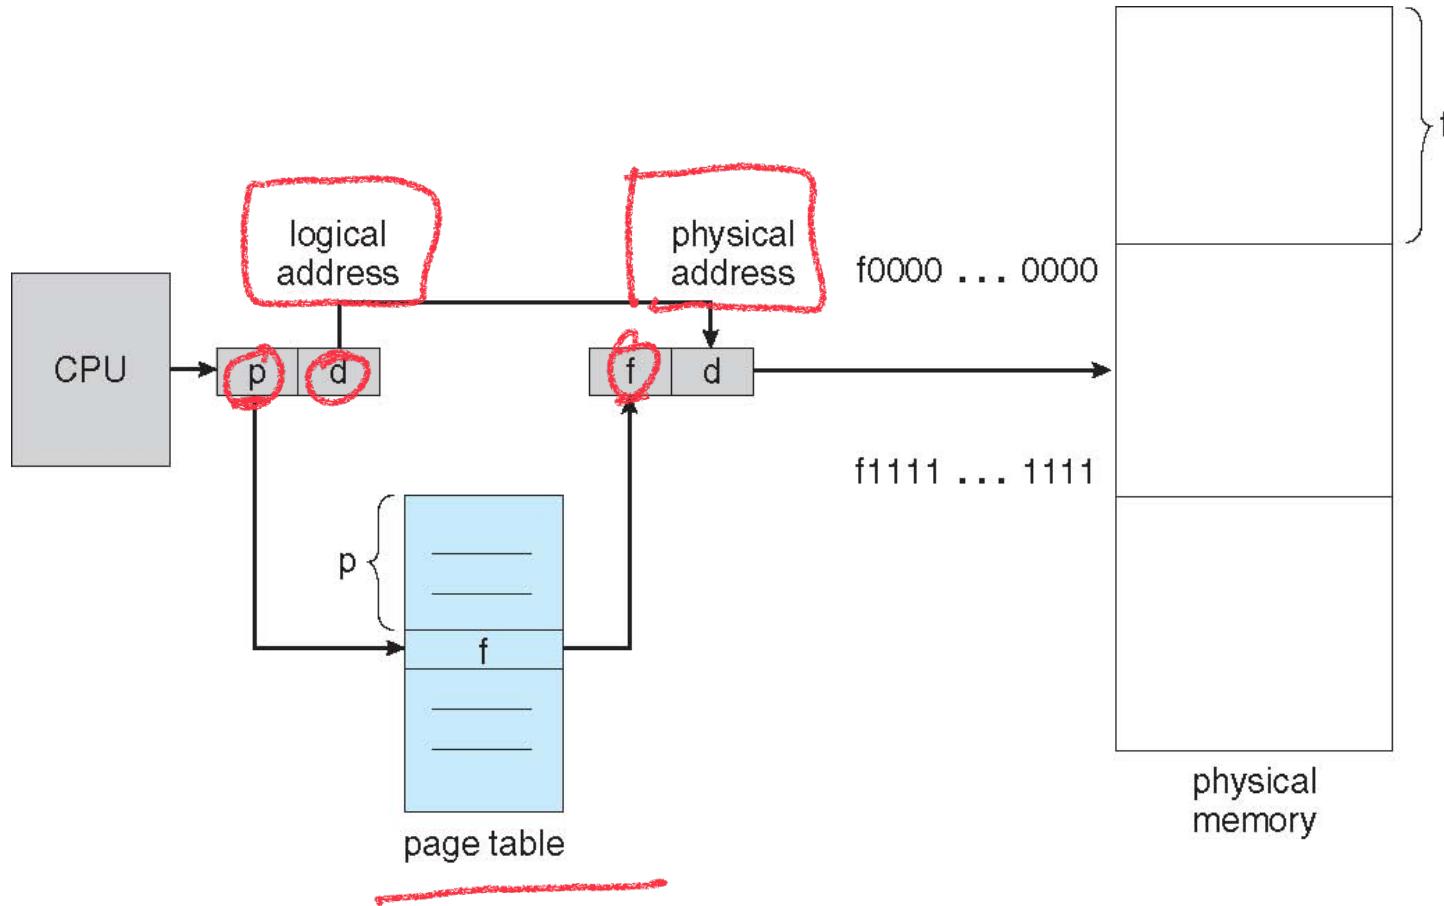


- For given logical address space 2^m
 - page size $2^n \rightarrow$ page offset → determine the size of the page / few
 - page numbers $2^{m-n} \rightarrow$ determine # of entries in the page table





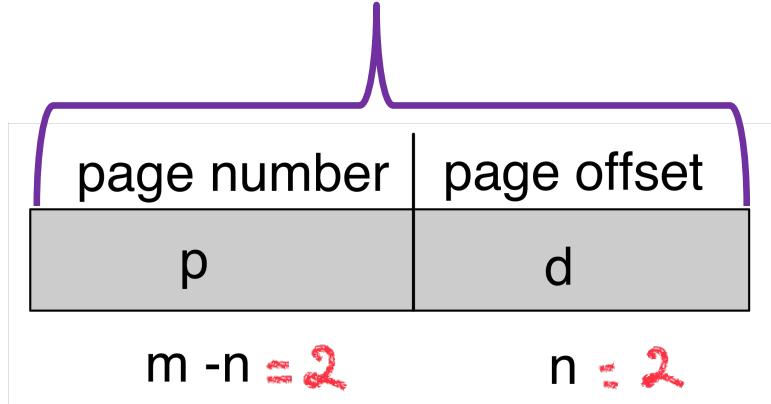
Paging Hardware





Example:

$m = \text{total number of bits} = 4$



Suppose:

$$m = 4, n = 2$$

Then:

$$\begin{aligned}\text{the total number of pages} &= 2^{m-n} = 2^{4-2} = 4 \quad 2^2 \\ \text{page size} &= 2^n = 2^2 = 4 \text{ Bytes}\end{aligned}$$





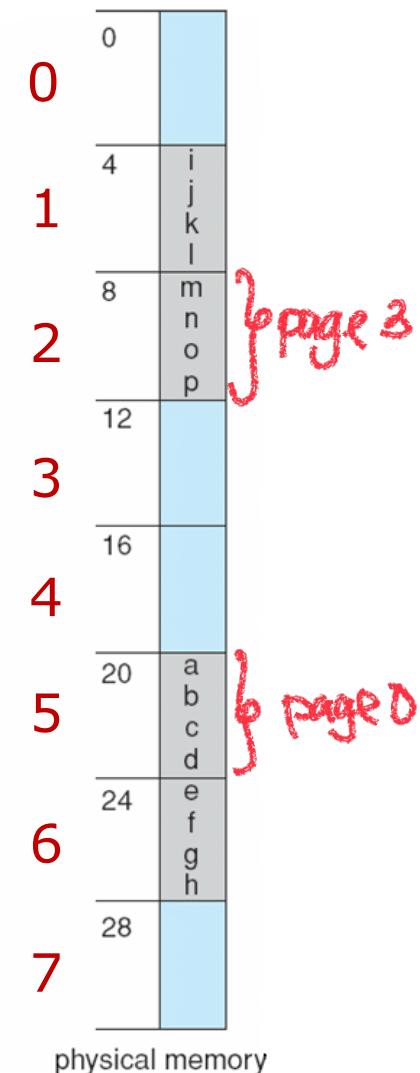
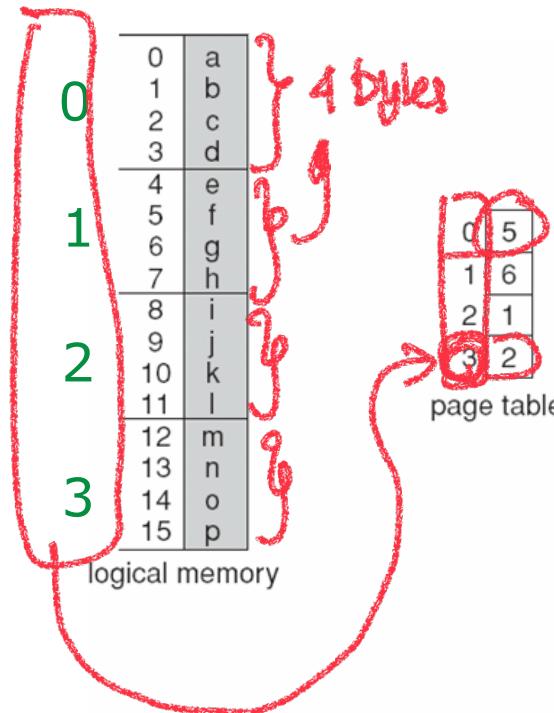
Paging Example

page number	page offset
p	d
m - n	n

$n=2$ and $m=4$

Total number of pages = $2^{m-n} = 2^2 = 4$

Each page contains $2^n = 2^2 = 4$ bytes
(page size)



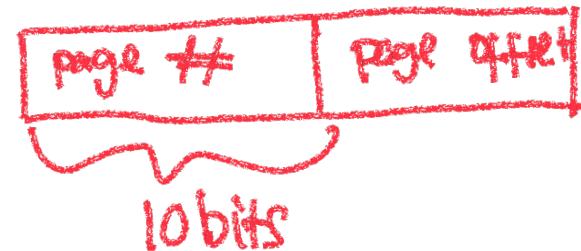


Check your understanding

1

A **page number** is represented by 10 bits. How many entries are in the page table?

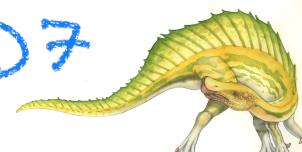
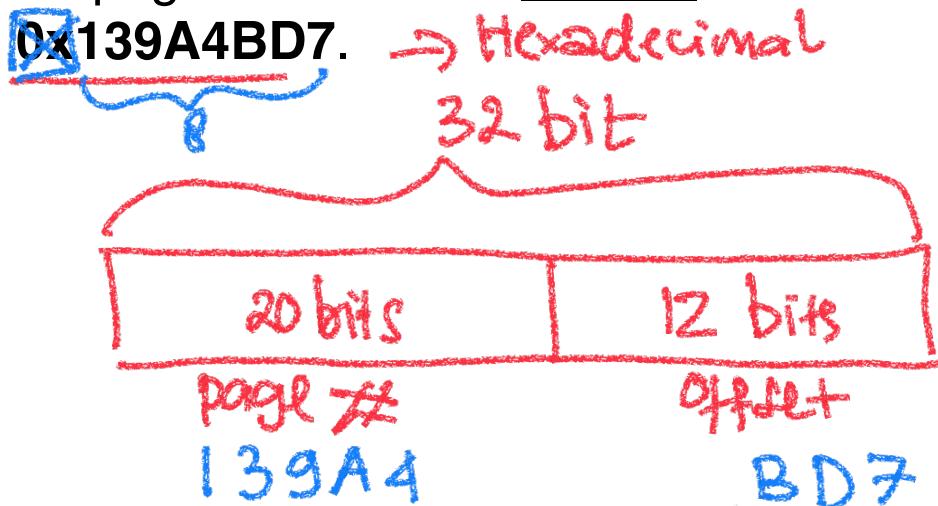
- A) 10
- B) 2^{10}



2

A 32-bit address is divided into 20 bits for the page number and 12 bits for the offset. _____ is the page number and _____ is the offset for the logical address $0x139A4BD7$. → Hexadecimal

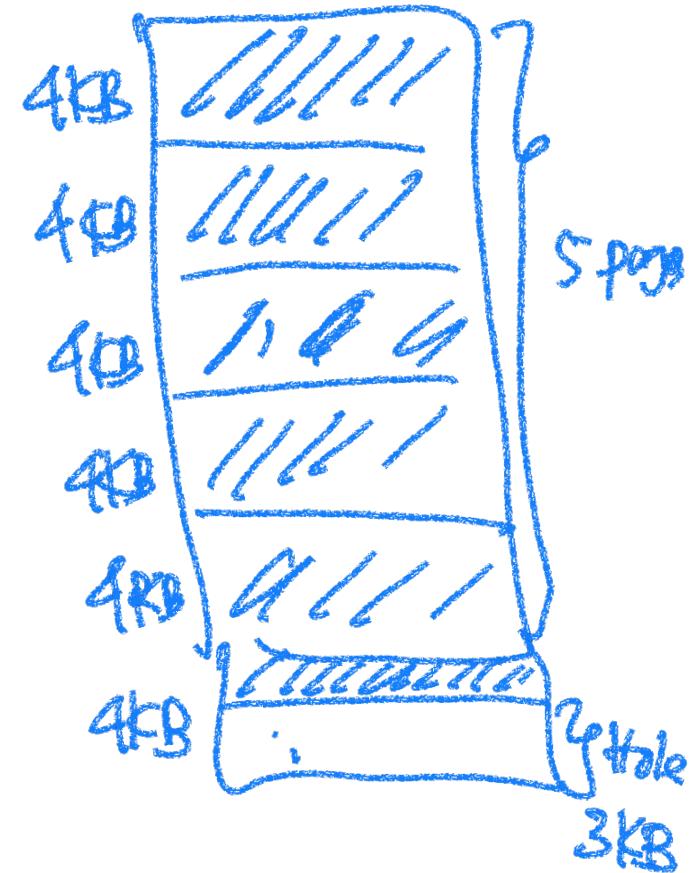
- A) 0x139, 0xA4BD7
- B) 0x139A, 0x4BD7
- C) 0x139A4, 0xBD7





21 KB

$$21/4 = 5.25$$





Paging (Cont.)

Calculating internal fragmentation

- Page size = 2,048 bytes = 2KB
- Process size = 72,766 bytes
- #pages required by the process = 36 pages
 - 72,766 / 2,048 = 35 pages + 1,086 bytes = 36 pages
 - 35 * 2,048 Bytes + 1,086 Bytes = 72,766 Bytes
- Internal fragmentation of 2,048 - 1,086 = 962 bytes → page 26
- Worst case fragmentation = 1 frame – 1 byte
- On average fragmentation = 1/2 frame size
- So small frame sizes desirable?
- But each page table entry takes memory to track
- Page sizes growing over time (typically 4 KB or 8 KB in size)
 - Solaris supports two page sizes – 8 KB and 4 MB
 - Windows 10 supports page sizes of 4 KB and 2 MB





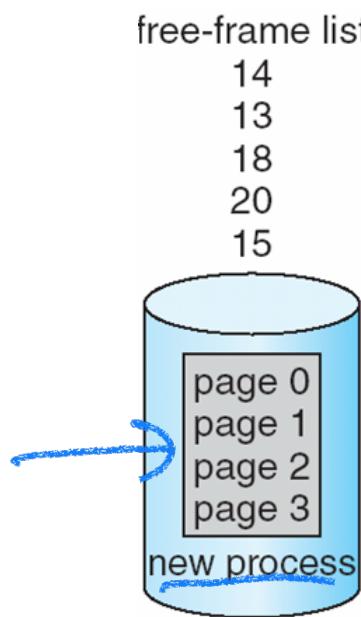
Paging (Cont.)

- Process view and physical memory now very different
 - An important aspect of paging is the clear separation between the programmer's view of memory and the actual physical memory.
 - In fact, the user program is scattered throughout physical memory, which also holds other programs.
- By implementation process can only access its own memory

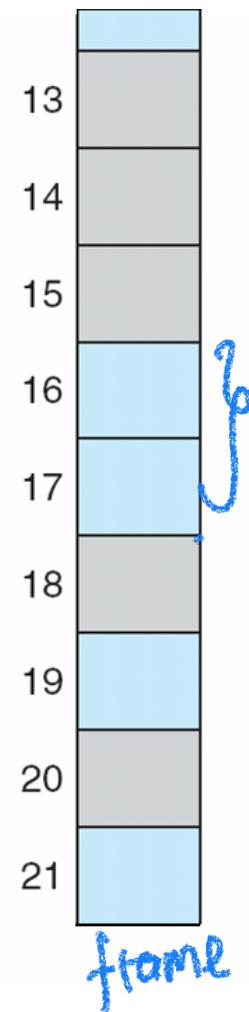




Free Frames

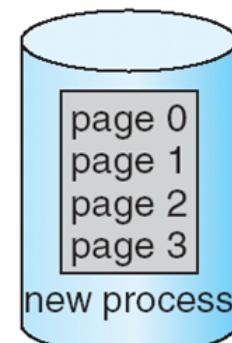


(a)



free-frame list

15



0	14
1	13
2	18
3	20

new-process page table

(b)

Before allocation

After allocation





Implementation of Page Table

- Page table is kept in main memory
- **Page-table base register (PTBR)** points to the page table
- **Page-table length register (PTLR)** indicates size of the page table
- In this scheme every data/instruction access requires two memory accesses
 - One for the page table and one for the data / instruction
- The two memory access problem can be solved by the use of a special fast-lookup hardware cache called **associative memory** or **translation look-aside buffers (TLBs)**





Implementation of Page Table (Cont.)

- Some TLBs store **address-space identifiers (ASIDs)** in each TLB entry – uniquely identifies each process to provide address-space protection for that process
 - Otherwise need to flush at every context switch
- TLBs typically small (64 to 1,024 entries)
- On a TLB miss, value is loaded into the TLB for faster access next time
 - Replacement policies must be considered
 - Used LRU (least recently used), round robin, or random
 - Some entries can be **wired down** for permanent fast access





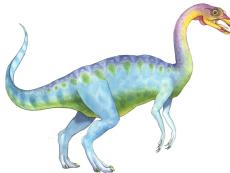
Associative Memory (TLBs)

- Associative memory – parallel search

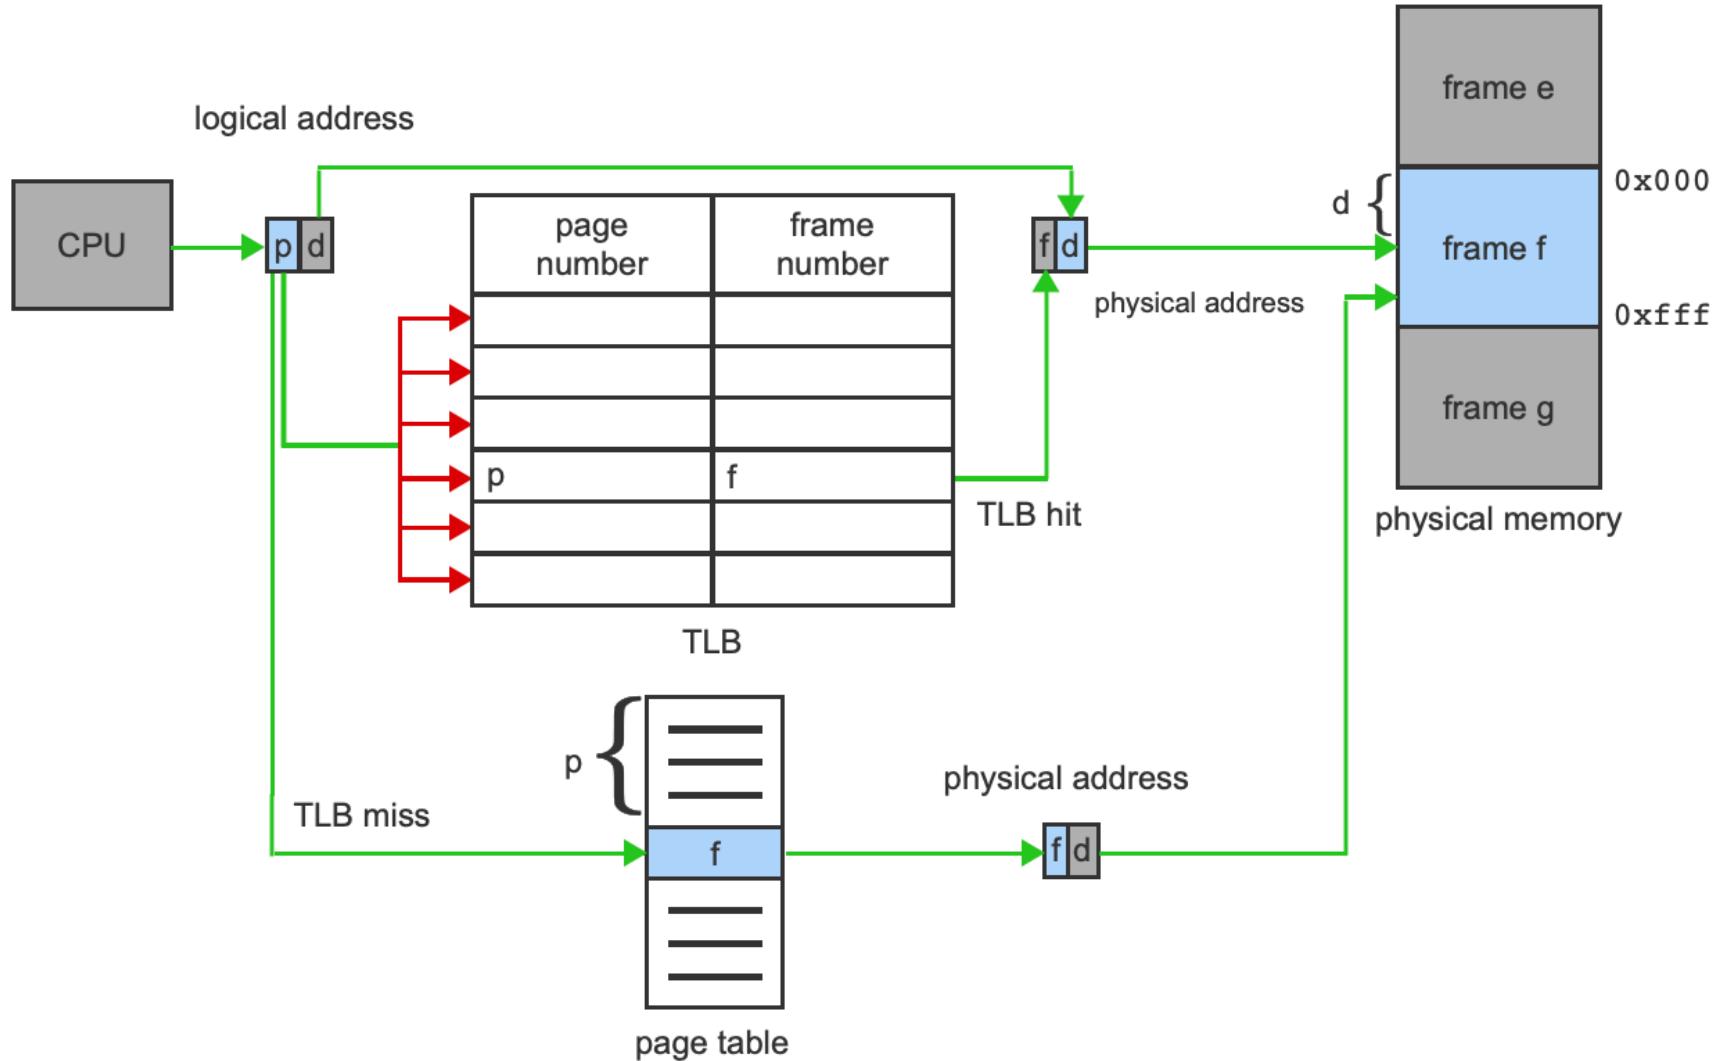
Page #	Frame #

- Address translation (p, d)
 - If p is in associative register, get frame # out
 - Otherwise get frame # from page table in memory





Paging Hardware With TLB





Effective Access Time

- Associative Lookup = ε time unit
 - Can be < 10% of memory access time
- Hit ratio = α
 - Hit ratio – percentage of times that a page number is found in the associative registers; ratio related to number of associative register
- **Effective Access Time (EAT)**
 $EAT = (1 + \varepsilon) \alpha + (2 + \varepsilon)(1 - \alpha)$
(when ε is omitted)

$$EAT = \alpha \times (1 \times \text{memory access time}) + (1 - \alpha) \times (2 \times \text{memory access time})$$

- Consider $\alpha = 80\%$, $\varepsilon = 20\text{ns}$ for TLB search, 100ns for memory access
 - $EAT = 0.80 \times 100 + 0.20 \times 200 = 120\text{ns}$ (ε is omitted)
- Consider more realistic hit ratio $\rightarrow \alpha = 99\%$, $\varepsilon = 20\text{ns}$ for TLB search, 100ns for memory access
 - $EAT = 0.99 \times 100 + 0.01 \times 200 = 101\text{ns}$





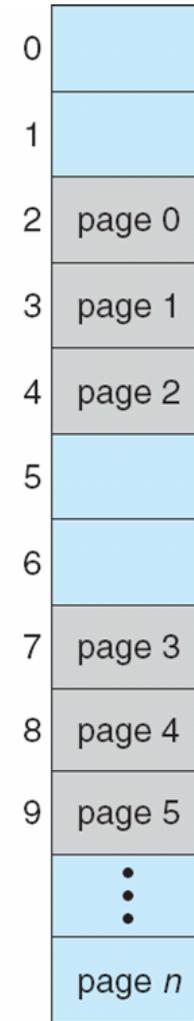
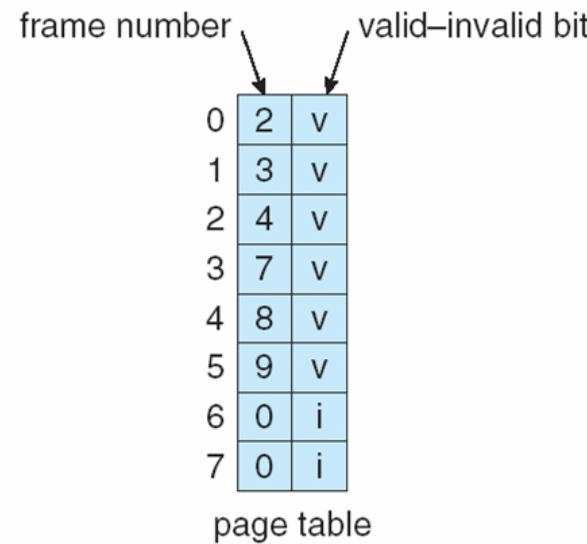
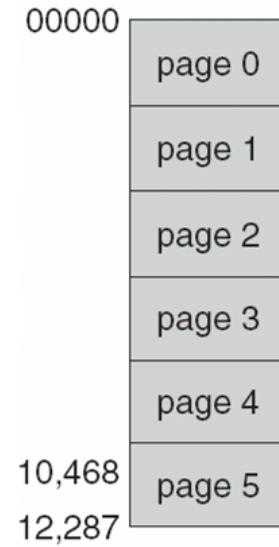
Memory Protection

- Memory protection implemented by associating protection bit with each frame to indicate if read-only or read-write access is allowed
 - Can also add more bits to indicate page execute-only, and so on
- **Valid-invalid** bit attached to each entry in the page table:
 - “valid” indicates that the associated page is in the process’ logical address space, and is thus a legal page
 - “invalid” indicates that the page is not in the process’ logical address space
 - Or use **page-table length register (PTLR)**
- Any violations result in a trap to the kernel





Valid (v) or Invalid (i) Bit In A Page Table





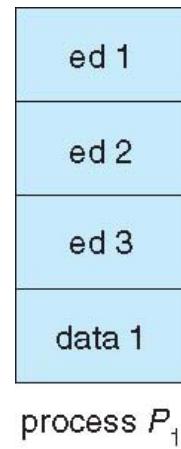
Shared Pages

- **Shared code**
 - One copy of read-only (**reentrant**) code shared among processes (i.e., text editors, compilers, window systems)
 - Similar to multiple threads sharing the same process space
 - Also useful for interprocess communication if sharing of read-write pages is allowed
- **Private code and data**
 - Each process keeps a separate copy of the code and data
 - The pages for the private code and data can appear anywhere in the logical address space



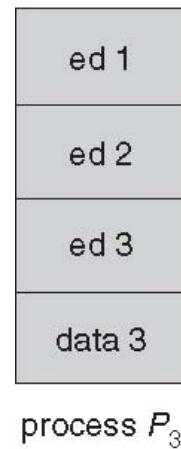


Shared Pages Example



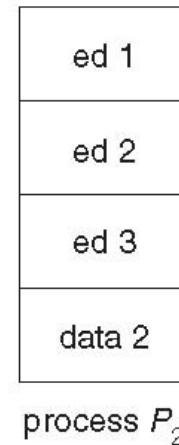
page table
for P_1

3
4
6
1



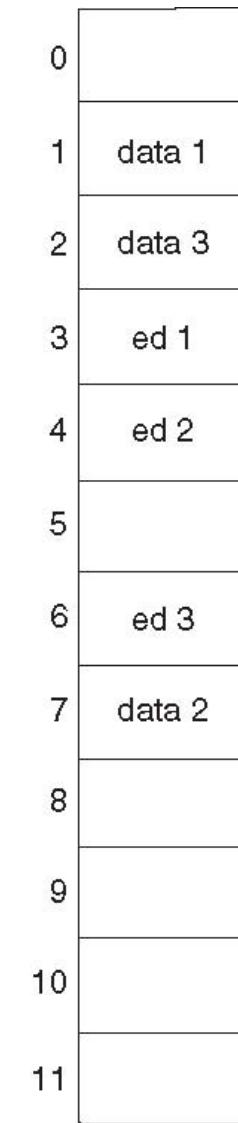
page table
for P_3

3
4
6
2



page table
for P_2

3
4
6
7





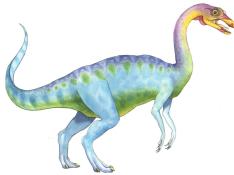
Check your understanding

- 1) A logical address space of size 2^{32} with a page size of 2^{12} requires _____ bits for the page number and _____ bits for the page offset.
 - A) 12, 20
 - B) 20, 12
 - C) 32, 12

- 2) How many bits must be used to represent the page offset for a page size of 8 KB?
 - A) 8
 - B) 10
 - C) 13

- 3) Paging can have _____ fragmentation but not _____
 - A) internal, external
 - B) external, internal



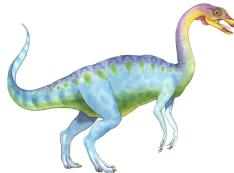


Check your understanding

- 6) What is the effective memory access time with a TLB hit ratio of 85% and 15 ns memory access time?

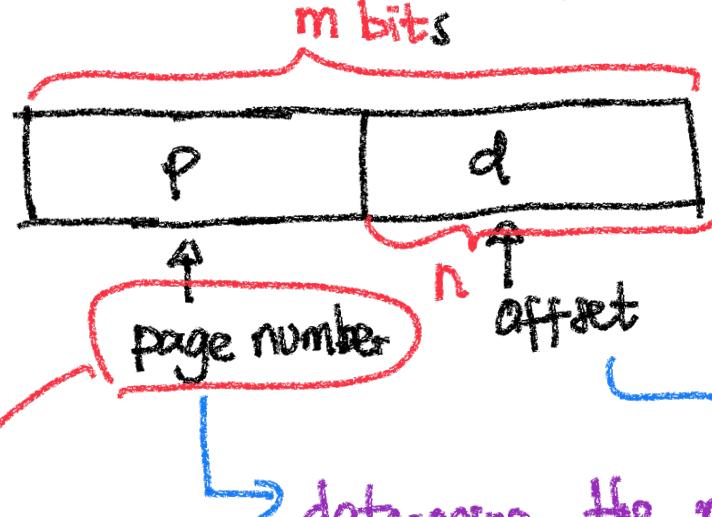
- 7) What is the page offset of the logical address 0xAF9 with a page size of 256 bytes?





Paging

Recall the structure of logical address in paging



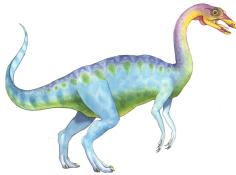
→ determine the size of each page or frame

→ determine the number of entries in the page table

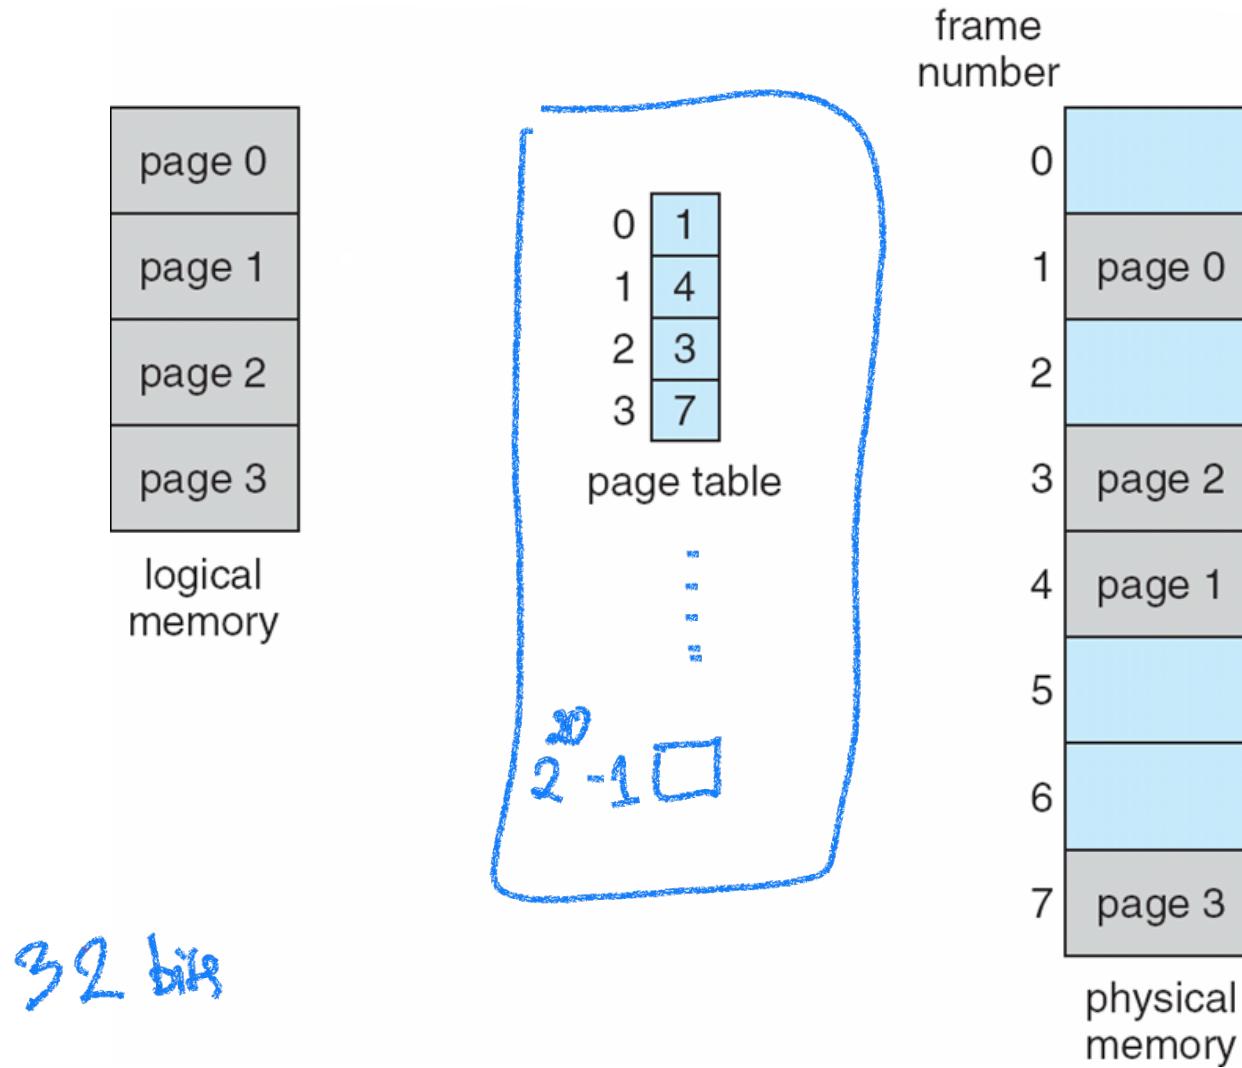
If offset d is n bits → the size of each page = ? 2^n
1 address \approx 1 byte

$m-n$ bits \Rightarrow # entries in page table = 2^{m-n}





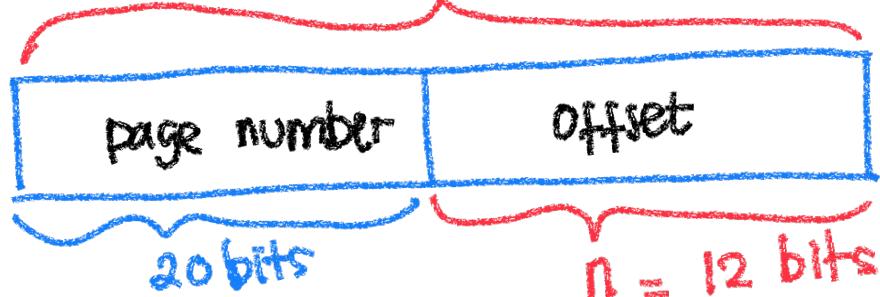
Paging Model of Logical and Physical Memory





$$1\text{ KB} = 2^{10}\text{ B}$$

32 bits logical address
 $m = 32$



Page size = 4KB \approx offset

$$n \Rightarrow \underline{4\text{KB}} = \underline{2^2} \times \underline{2^{10}} \text{ Bytes} = 2^{12} \text{ Bytes}$$

$n = 12$

How many bits used for page number? $32 - 12 = 20$

What is the size of the page table? $2^{20} \text{ Bytes} = 1 \text{ MB}$





Structure of the Page Table

- Memory structures for paging can get huge using straight-forward methods
 - Consider a 32-bit logical address space as on modern computers
 - Page size of 4 KB (2^{12}) → $4 \times 2^{10} \text{ Bytes} = 2^2 \times 2^{10} \text{ Bytes}$
 - Page table would have 1 million entries ($2^{32} / 2^{12}$)
 - If each entry is 4 bytes -> 4 MB of physical address space / memory for page table alone
 - That amount of memory used to cost a lot
 - Don't want to allocate that contiguously in main memory
- Hierarchical Paging
- Hashed Page Tables
- Inverted Page Tables

$$1 \text{ KB} = 2^{10} \text{ Bytes}$$

$$1 \text{ MB} = 2^{20} \text{ Bytes}$$

$$1 \text{ GB} = 2^{30} \text{ Bytes}$$





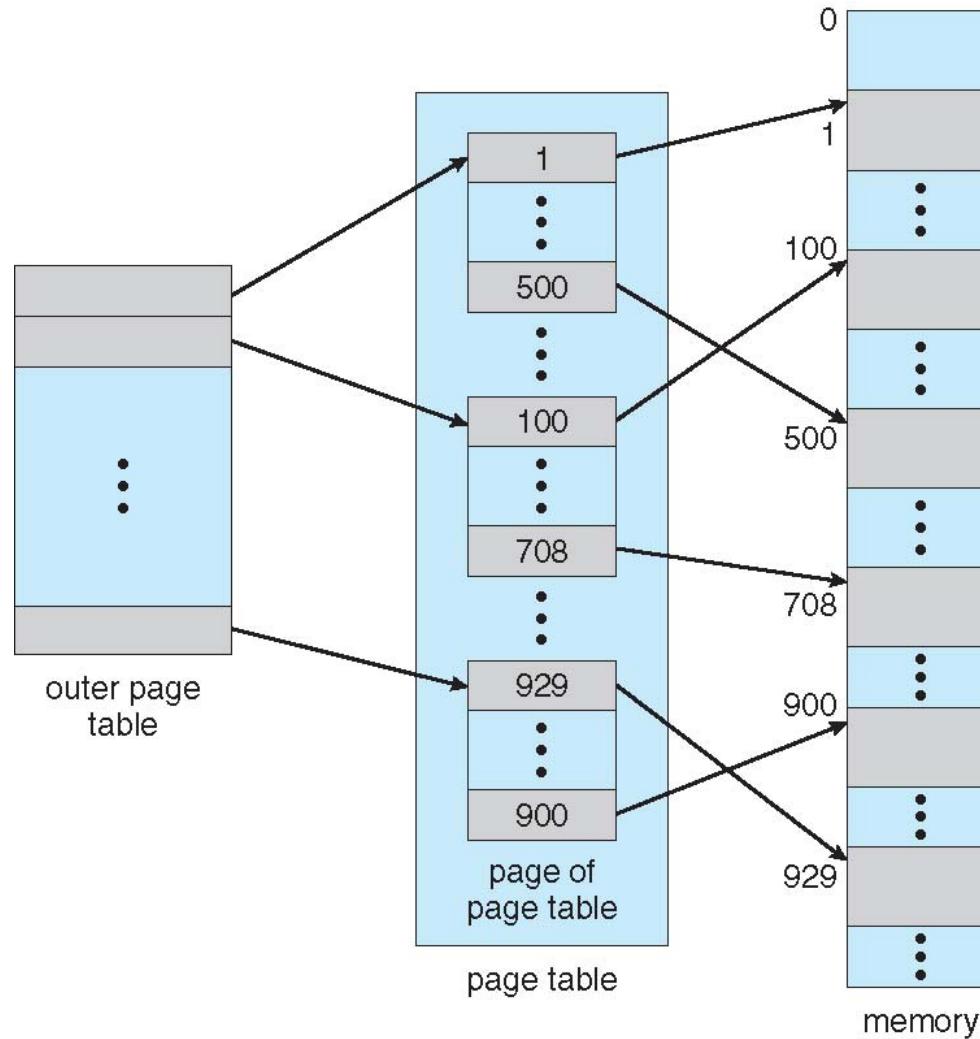
Hierarchical Page Tables

- Break up the logical address space into multiple page tables
- A simple technique is a two-level page table
- We then page the page table





Two-Level Page-Table Scheme



address translation works from the outer page table inward,
also known as a ***forward-mapped*** page table.

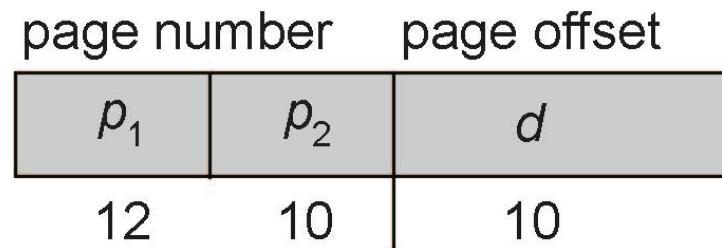




Two-Level Paging Example

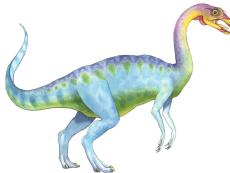
- A logical address (on 32-bit machine with 1K page size) is divided into:
 - a page number consisting of 22 bits
 - a page offset consisting of 10 bits
- Since the page table is paged, the page number is further divided into:
 - a 12-bit page number
 - a 10-bit page offset
- Thus, a logical address is as follows:

→ the size of each page

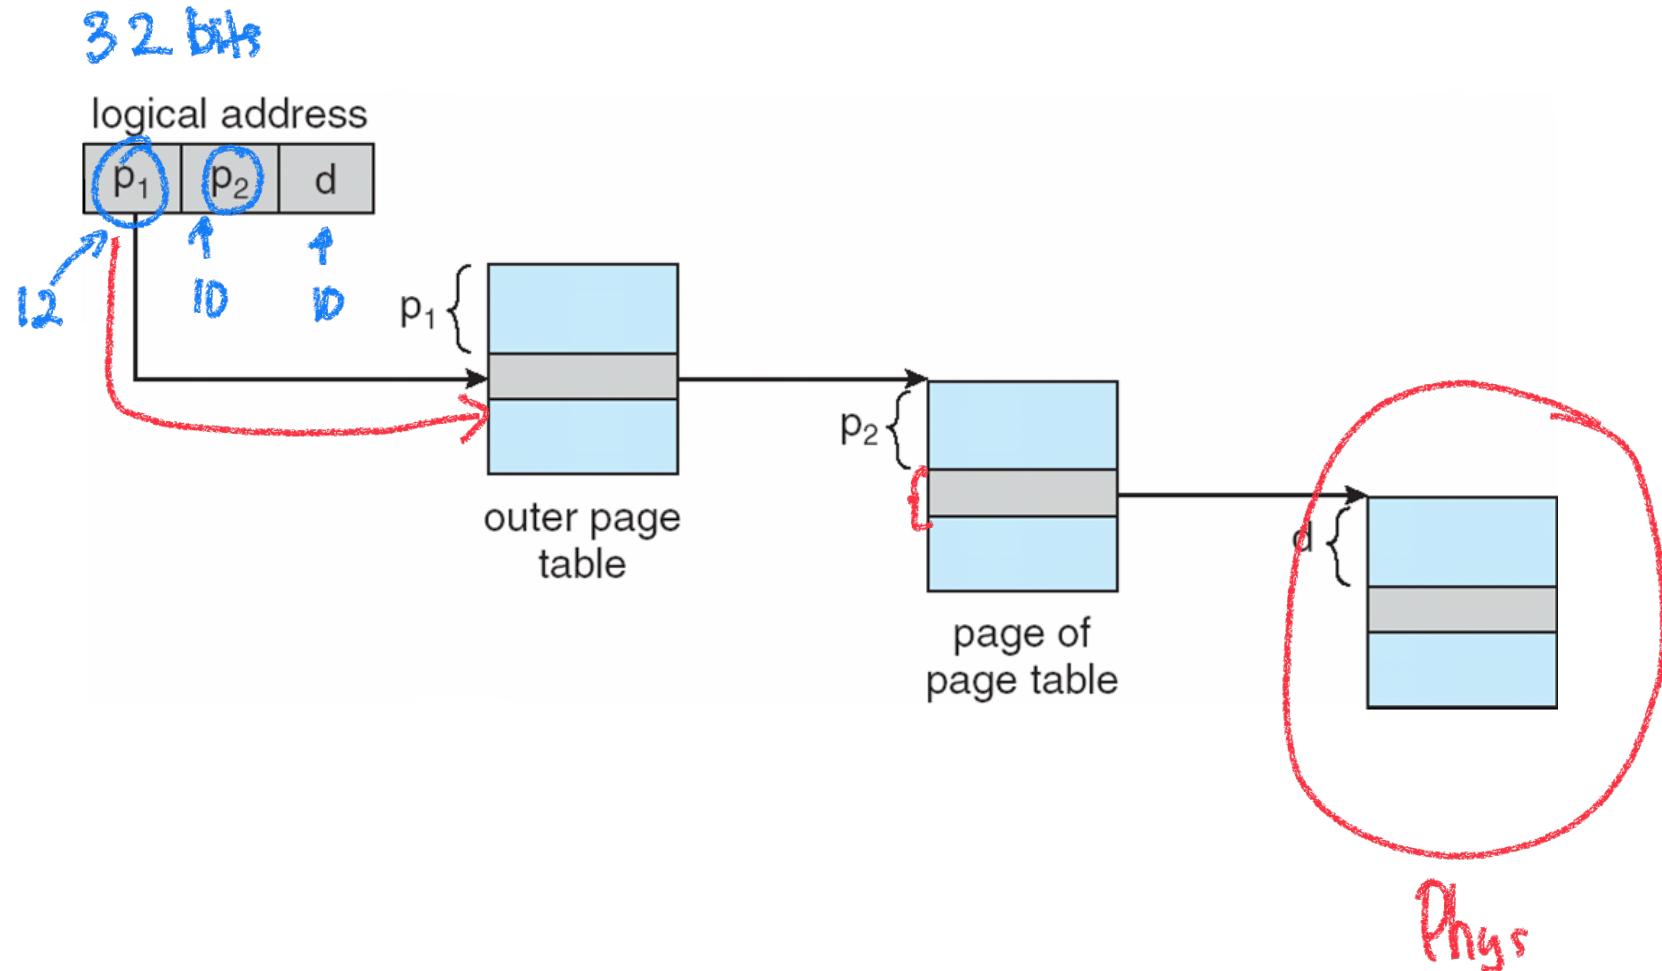


- where p_1 is an index into the outer page table, and p_2 is the displacement within the page of the inner page table
- Known as **forward-mapped page table**





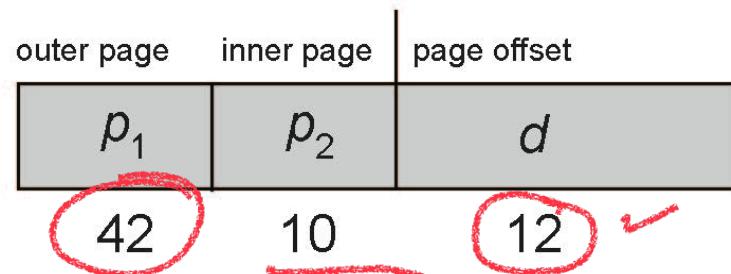
Address-Translation Scheme





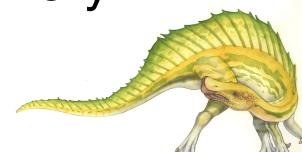
64-bit Logical Address Space

- Even two-level paging scheme not sufficient
- If page size is 4 KB (2^{12})
 - Then page table has 2^{52} entries
 - If two level scheme, inner page tables could be 2^{10} 4-byte entries
 - Address would look like



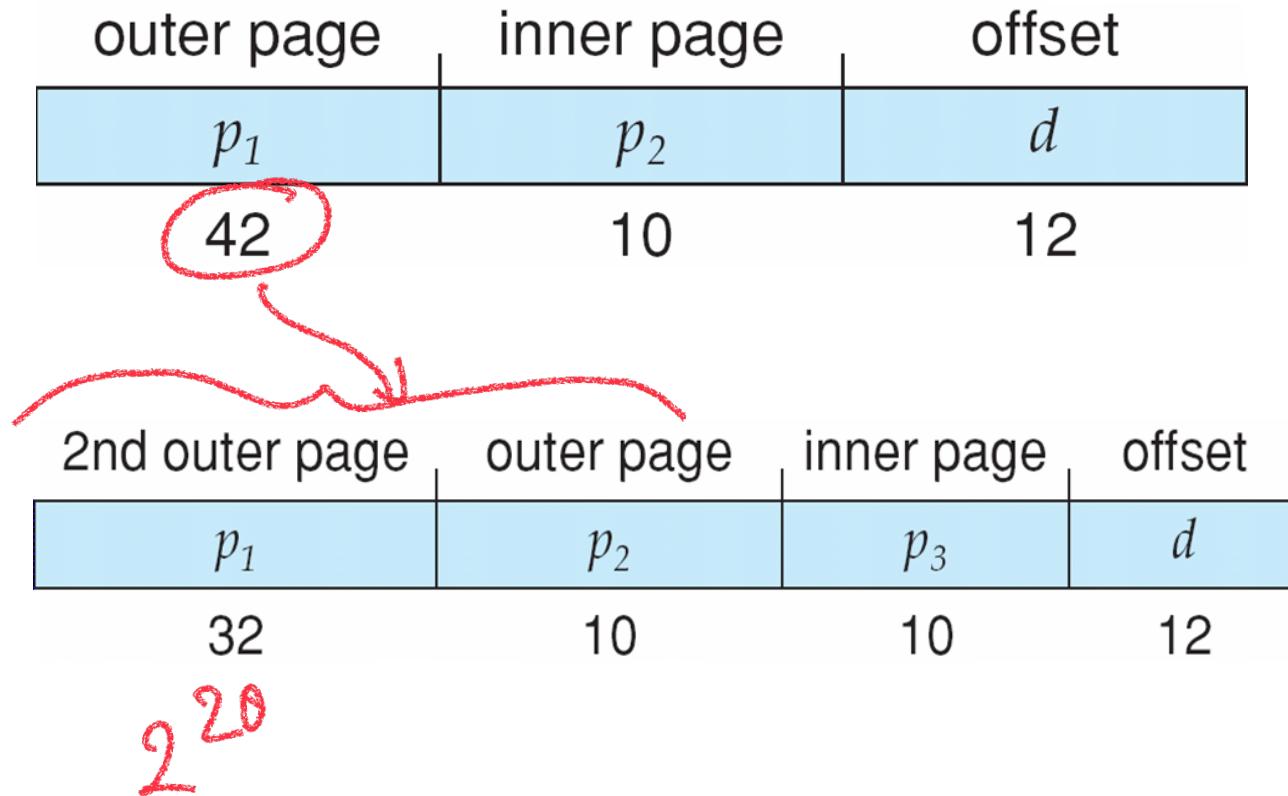
2^{42}

- Outer page table has 2^{42} entries or 2^{44} bytes
- One solution is to add a 2nd outer page table
- But in the following example the 2nd outer page table is still 2^{34} bytes in size
 - And possibly 4 memory access to get to one physical memory location





Three-level Paging Scheme





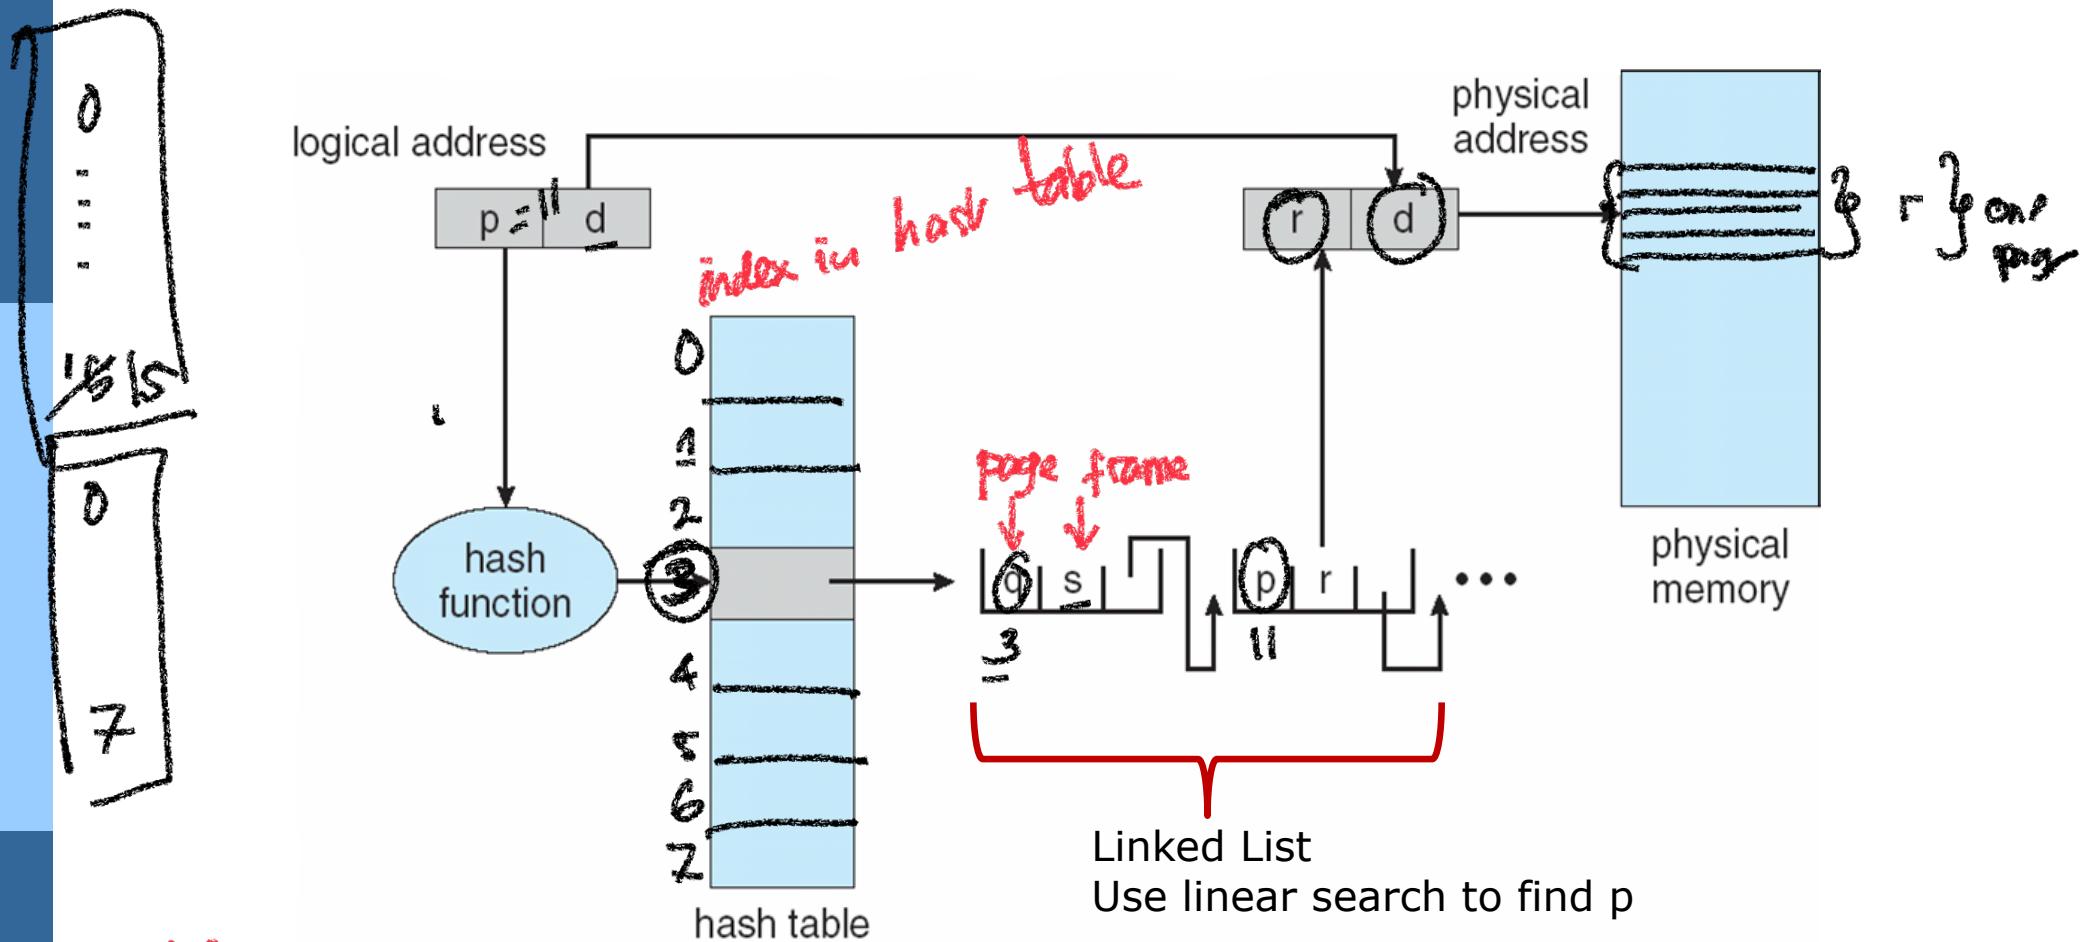
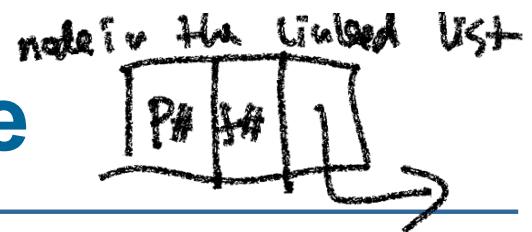
Hashed Page Tables

- Common in address spaces > 32 bits
- The virtual page number is hashed into a page table
 - This page table contains a chain of elements hashing to the same location
- Each element contains (1) the virtual page number (2) the value of the mapped page frame (3) a pointer to the next element
- Virtual page numbers are compared in this chain searching for a match
 - If a match is found, the corresponding physical frame is extracted
- Variation for 64-bit addresses is **clustered page tables**
 - Similar to hashed but each entry refers to several pages (such as 16) rather than 1
 - Especially useful for **sparse** address spaces (where memory references are non-contiguous and scattered)





Hashed Page Table



e.g.

$M = 8 \text{ bits}$

$n = 4 \text{ bits offset} \rightarrow \text{page size} = 2^4 = 16 \text{ Bytes}$

$$\text{hash}(p) = p \bmod 8$$

$$= 3 \bmod 8 : 3$$

$2^4 \rightarrow \# \text{ entries in page table}$





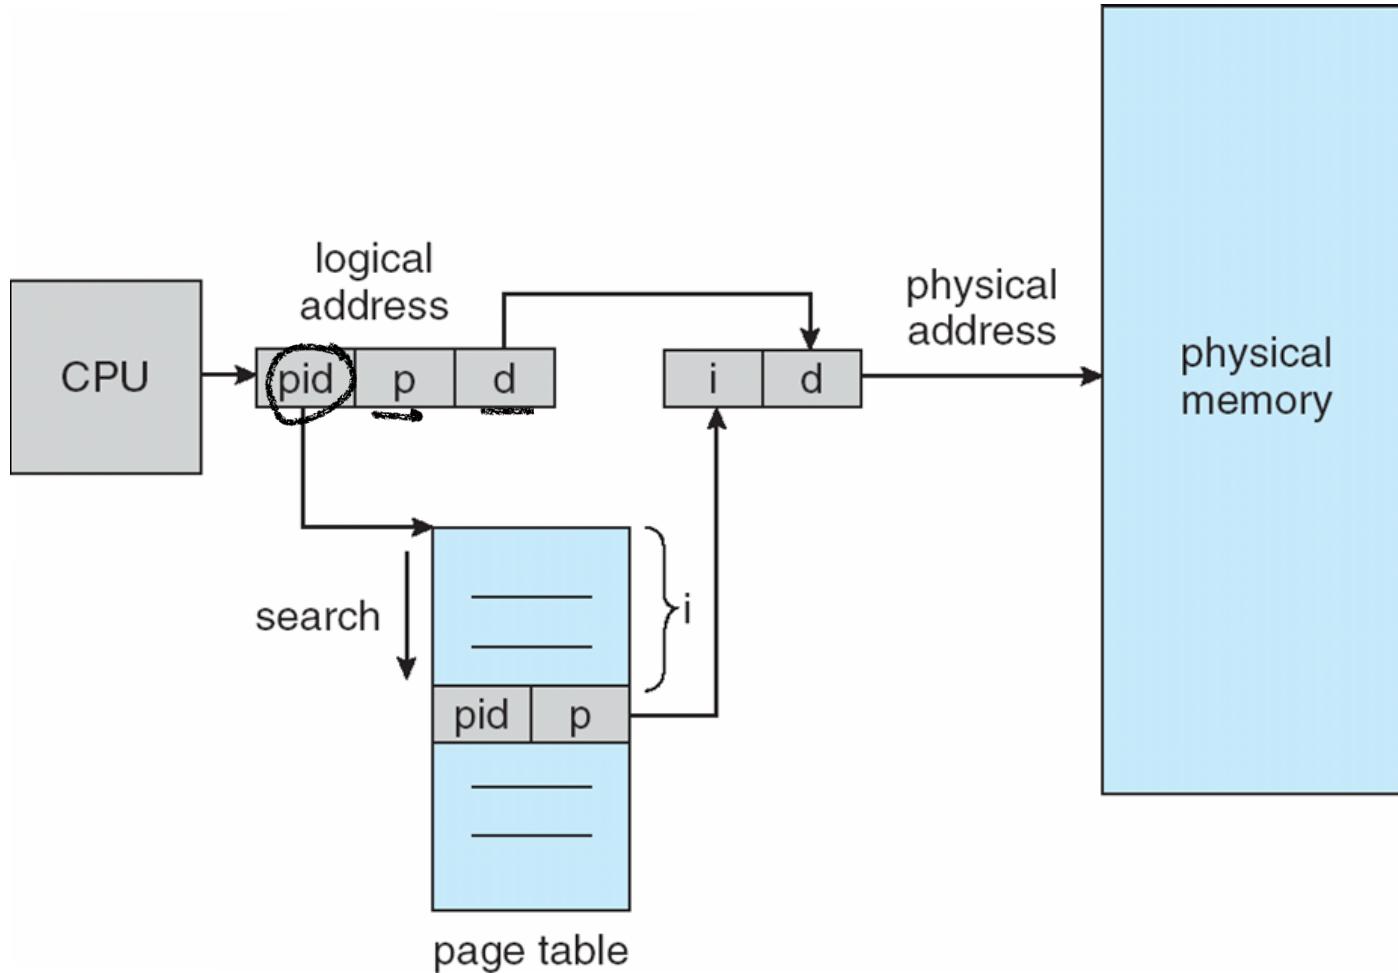
Inverted Page Table

- Rather than each process having a page table and keeping track of all possible logical pages, track all physical pages
- One entry for each real page of memory
- Entry consists of the virtual address of the page stored in that real memory location, with information about the process that owns that page
- Decreases memory needed to store each page table, but increases time needed to search the table when a page reference occurs
- Use hash table to limit the search to one — or at most a few — page-table entries
 - TLB can accelerate access
- But how to implement shared memory?
 - One mapping of a virtual address to the shared physical address



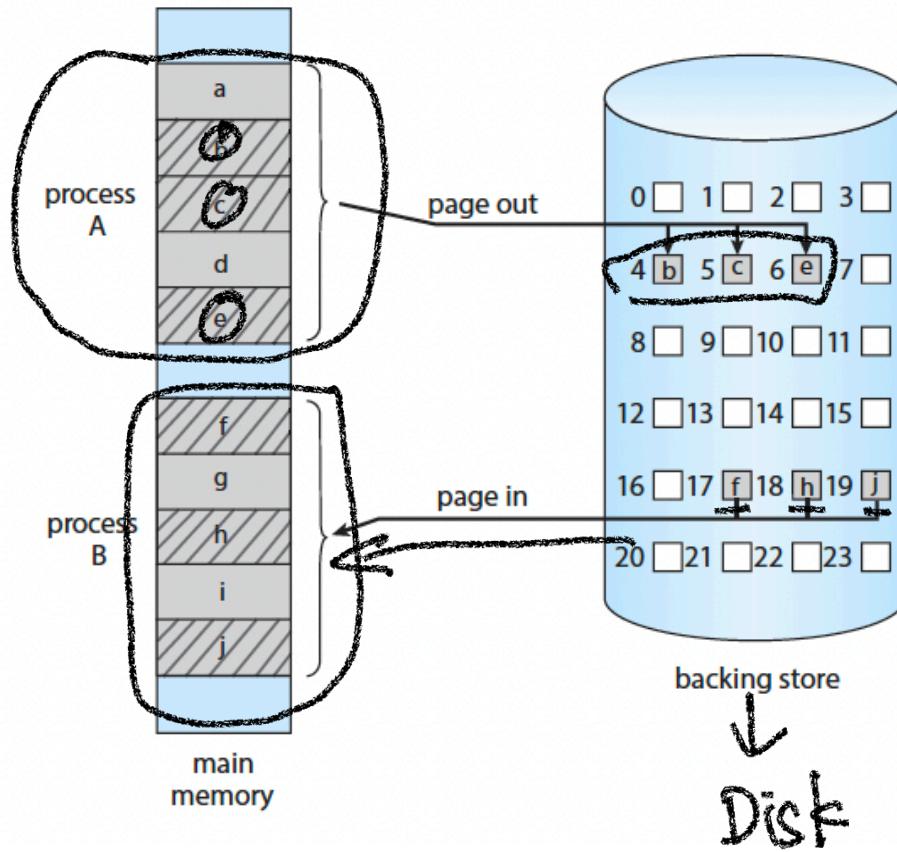


Inverted Page Table Architecture





Swapping with Paging

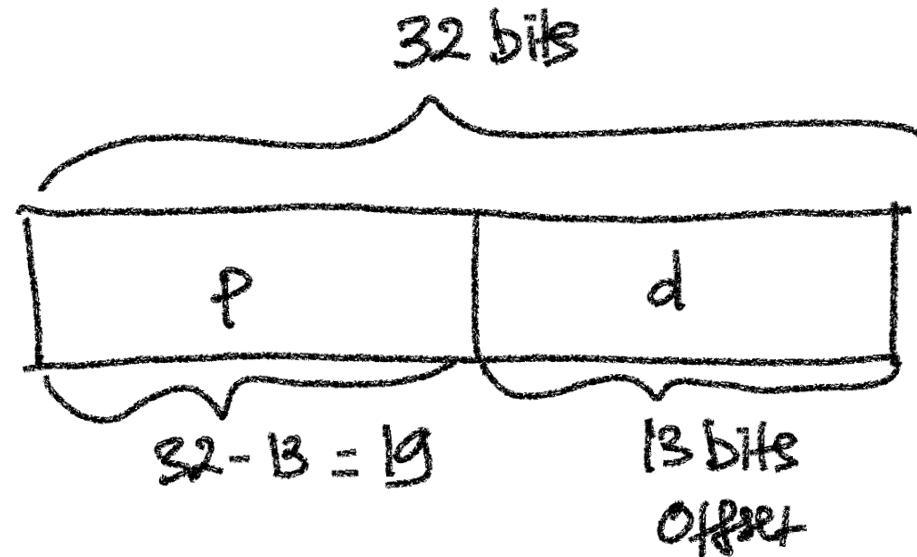




$$1\text{ KB} = 2^{10}$$

Practice TopHat 9.1

A 32-bit single-level paging system with an 8 KB page size has _____ entries in the page table.



$$\begin{aligned} 8\text{ KB} &= 2^? \text{ Bytes} \\ 2^3 \times 2^{10} \text{ Bytes} &= \\ 2^{3+10} &= 2^{13} \end{aligned}$$

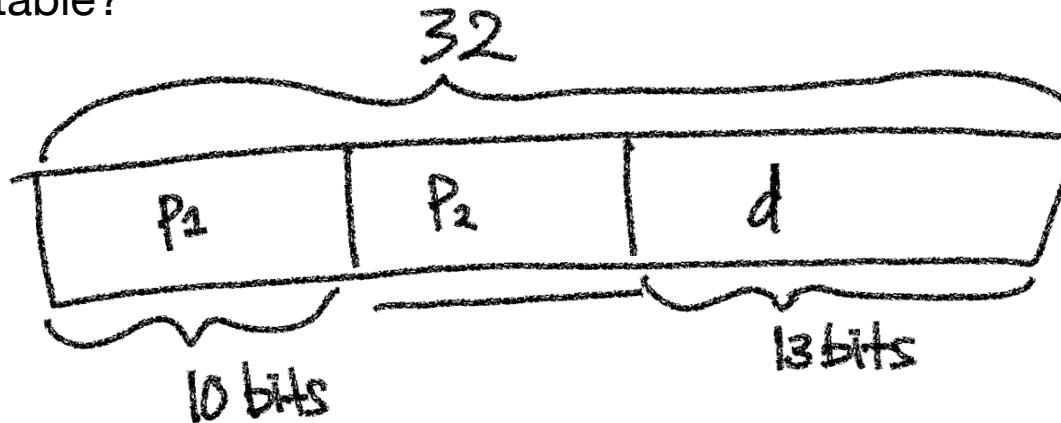
$$\begin{aligned} 2^0 &= 1 \\ 2^1 &= 2 \\ 2^2 &= 4 \\ 2^3 &= 8 \end{aligned}$$





Practice TopHat 9.2

Consider a 32-bit address for a two-level paging system with an 8 KB page size. The outer page table has 1024 entries. How many bits are used to represent the second-level page table?



$$8\text{KB} = 2^{13} \text{ Bytes}$$

1024 entries : 1024 Bytes

$$1\text{KB} = 1024 \text{ Bytes}$$

$$2^{10} \text{ Bytes}$$

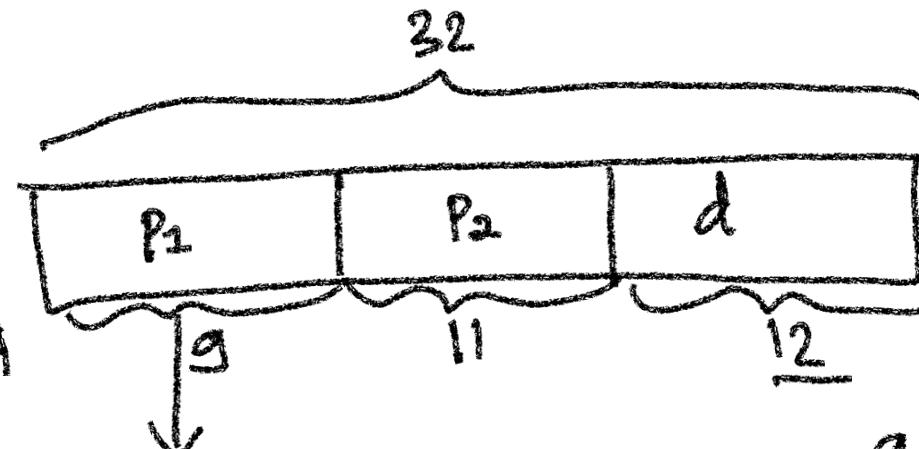




Practice TopHat 9.3

A two-level paging system with 32 bit addresses has 512 entries for the outer page table and 2048 entries for the inner page table. What is the size of a page?

$$2^{10} = \frac{1024 \times 2}{2}$$
$$2^9 = 512$$



$$1 \text{ KB} = 2^{10} = 1024$$

$$1 \text{ MB} = 2^{20}$$

$$1 \text{ GB} = 2^{30}$$

$$512 \text{ entries} = 512 \text{ B} \rightarrow 2^9 \text{ B}$$

$$2048 \text{ entries} = \frac{2048 \text{ B}}{4+2} \rightarrow 2^{11} \text{ B}$$

$$d = 12$$

$$\text{page size} = 2^{12} \text{ bytes} = \underline{\underline{4096}}$$





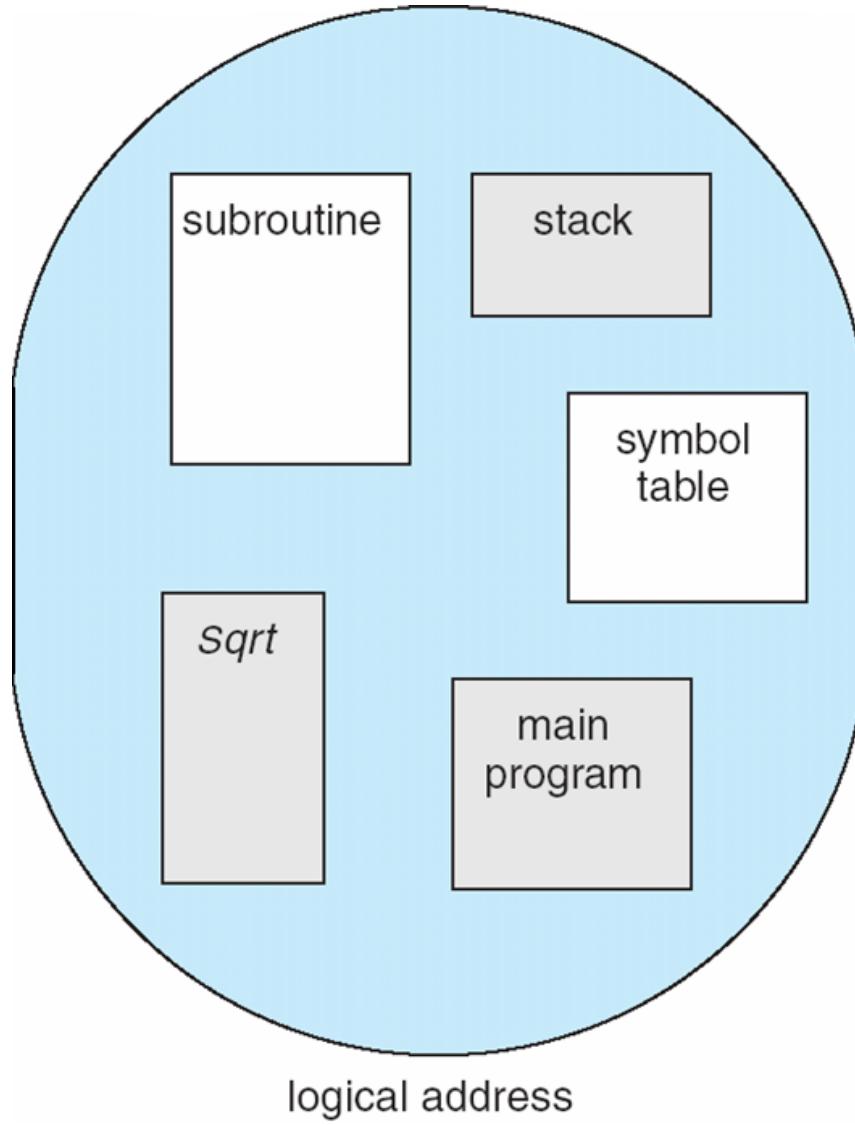
Segmentation

- Memory-management scheme that supports user view of memory
- A program is a collection of segments
 - A segment is a logical unit such as:
 - main program
 - procedure
 - function
 - method
 - object
 - local variables, global variables
 - common block
 - stack
 - symbol table
 - arrays



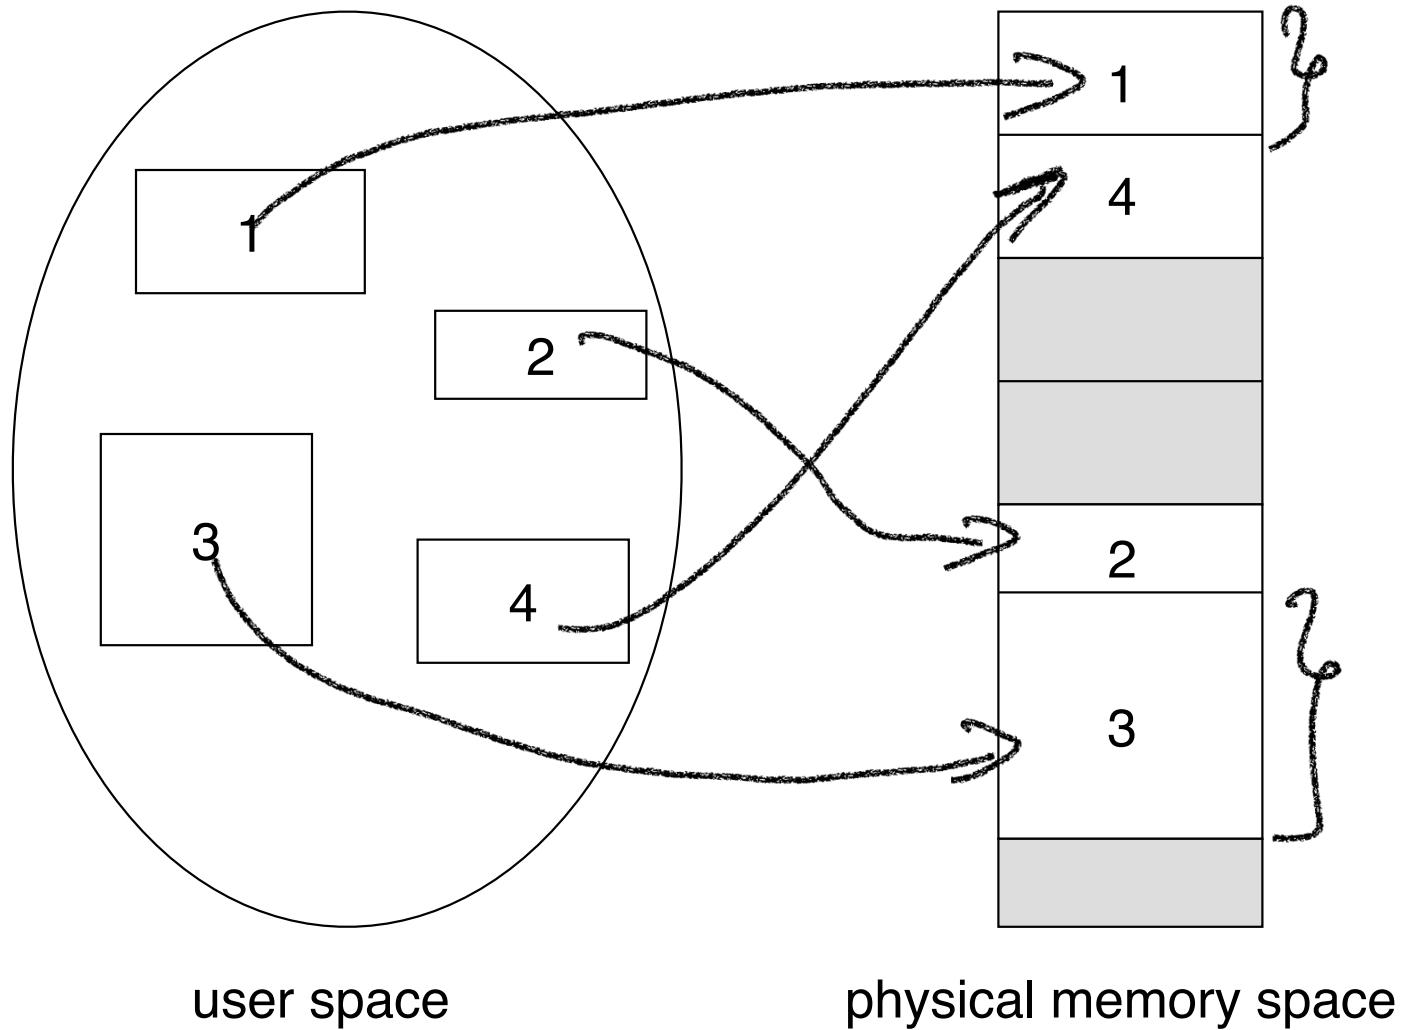


User's View of a Program





Logical View of Segmentation





Segmentation Architecture

- Logical address consists of a two tuple:
 $\langle \text{segment-number}, \text{offset} \rangle,$
- **Segment table** – maps two-dimensional physical addresses; each table entry has:
 - **base** – contains the starting physical address where the segments reside in memory
 - **limit** – specifies the length of the segment
- **Segment-table base register (STBR)** points to the segment table's location in memory
- **Segment-table length register (STLR)** indicates number of segments used by a program;
segment number **s** is legal if **$s < \text{STLR}$**





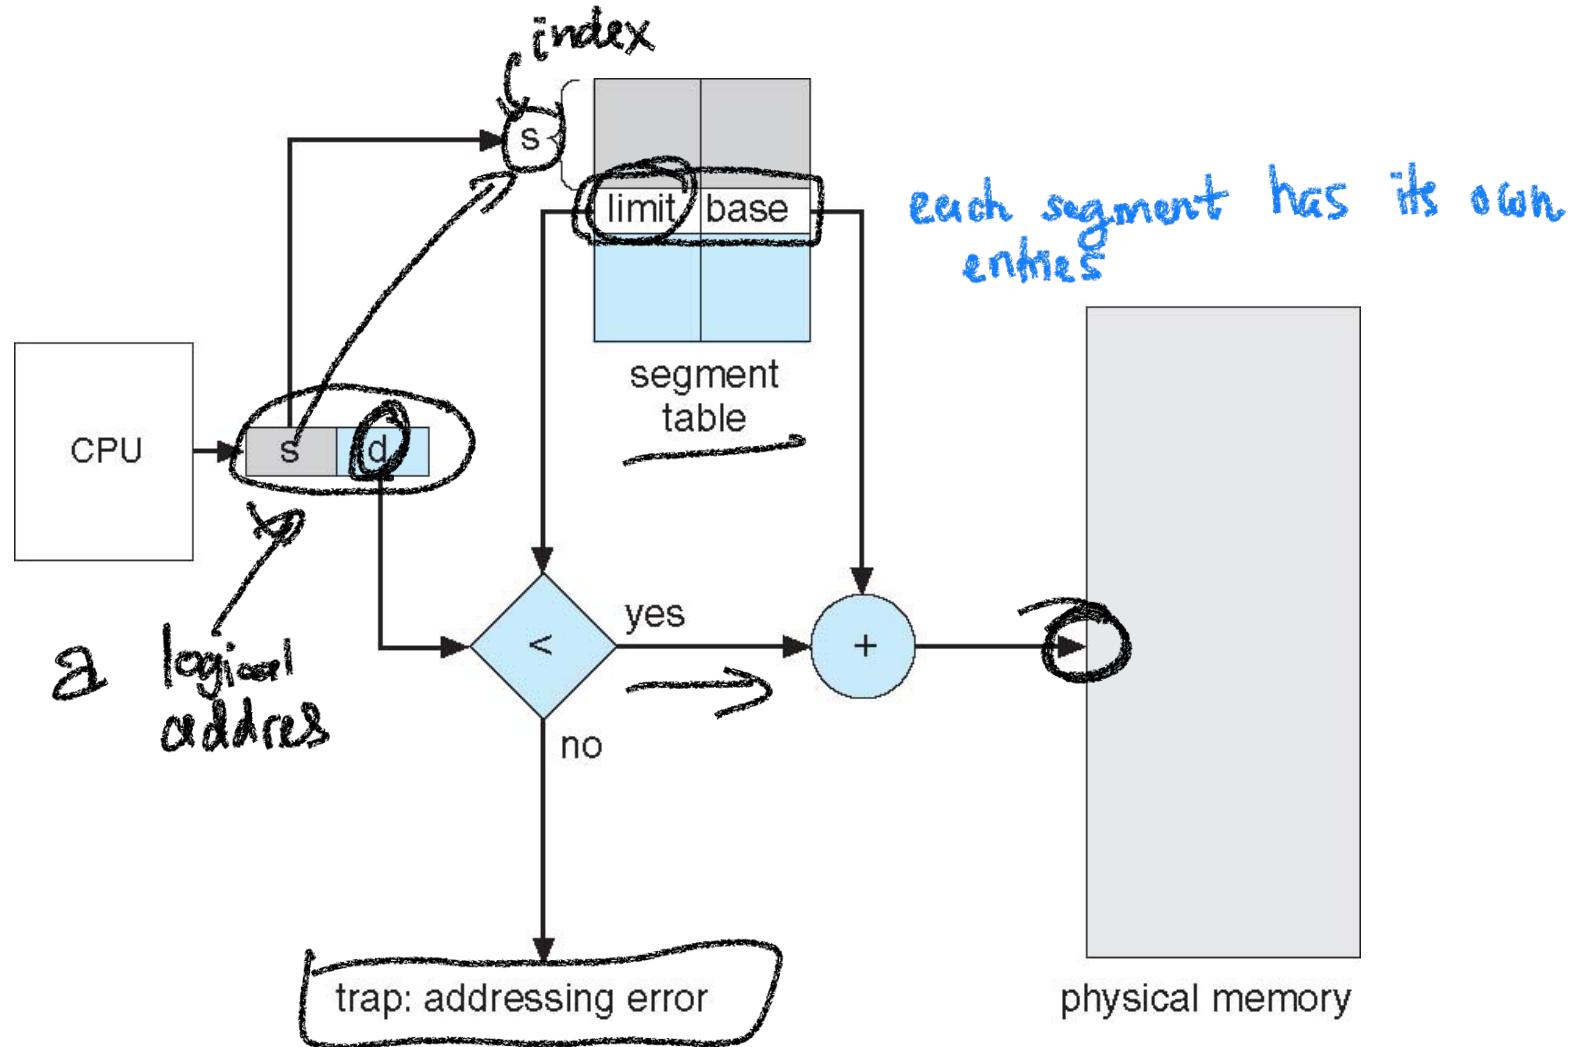
Segmentation Architecture (Cont.)

- Protection
 - With each entry in segment table associate:
 - validation bit = 0 \Rightarrow illegal segment
 - read/write/execute privileges
- Protection bits associated with segments; code sharing occurs at segment level
- Since segments vary in length, memory allocation is a dynamic storage-allocation problem
- A segmentation example is shown in the following diagram





Segmentation Hardware



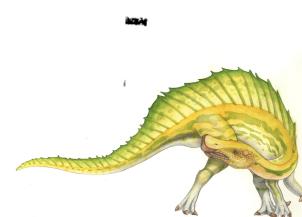
e.g.

s	d
---	---

$\langle 3 > 80 \rangle$

$s = 3$ in segment table

400	100
limit	base



End of Chapter 9

